

Choosing things in dataframes

Packages

The usual:

```
library(tidyverse)
```

Doing things with data frames

Let's go back to our Australian athletes:

```
athletes
```

```
# A tibble: 202 x 13
```

	Sex	Sport	RCC	WCC	Hc	Hg	Ferr	BMI	SSF	~%Bfat~	LBM
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	female	Netba~	4.56	13.3	42.2	13.6	20	19.2	49	11.3	53.1
2	female	Netba~	4.15	6	38	12.7	59	21.2	110.	25.3	47.1
3	female	Netba~	4.16	7.6	37.5	12.3	22	21.4	89	19.4	53.4
4	female	Netba~	4.32	6.4	37.7	12.3	30	21.0	98.3	19.6	48.8
5	female	Netba~	4.06	5.8	38.7	12.8	78	21.8	122.	23.1	56.0
6	female	Netba~	4.12	6.1	36.6	11.8	21	21.4	90.4	16.9	56.4
7	female	Netba~	4.17	5	37.4	12.7	109	21.5	107.	21.3	53.1
8	female	Netba~	3.8	6.6	36.5	12.4	102	24.4	157.	26.6	54.4
9	female	Netba~	3.96	5.5	36.3	12.4	71	22.6	101.	17.9	56.0
10	female	Netba~	4.44	9.7	41.4	14.1	64	22.8	126.	25.0	51.6

```
# i 192 more rows  
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Choosing a column

```
athletes %>% select(Sport)
```

```
# A tibble: 202 x 1
```

```
  Sport
```

```
  <chr>
```

```
1 Netball
```

```
2 Netball
```

```
3 Netball
```

```
4 Netball
```

```
5 Netball
```

```
6 Netball
```

```
7 Netball
```

```
8 Netball
```

```
9 Netball
```

```
10 Netball
```

```
# i 192 more rows
```

Choosing several columns

```
athletes %>% select(Sport, Hg, BMI)
```

```
# A tibble: 202 x 3
```

	Sport	Hg	BMI
	<chr>	<dbl>	<dbl>
1	Netball	13.6	19.2
2	Netball	12.7	21.2
3	Netball	12.3	21.4
4	Netball	12.3	21.0
5	Netball	12.8	21.8
6	Netball	11.8	21.4
7	Netball	12.7	21.5
8	Netball	12.4	24.4
9	Netball	12.4	22.6
10	Netball	14.1	22.8

```
# i 192 more rows
```

Choosing consecutive columns

```
athletes %>% select(Sex:WCC)
```

```
# A tibble: 202 x 4
```

	Sex	Sport	RCC	WCC
	<chr>	<chr>	<dbl>	<dbl>
1	female	Netball	4.56	13.3
2	female	Netball	4.15	6
3	female	Netball	4.16	7.6
4	female	Netball	4.32	6.4
5	female	Netball	4.06	5.8
6	female	Netball	4.12	6.1
7	female	Netball	4.17	5
8	female	Netball	3.8	6.6
9	female	Netball	3.96	5.5
10	female	Netball	4.44	9.7

```
# i 192 more rows
```

Choosing all-but some columns

```
athletes %>% select(-(RCC:LBM))
```

```
# A tibble: 202 x 4
```

	Sex	Sport	Ht	Wt
	<chr>	<chr>	<dbl>	<dbl>
1	female	Netball	177.	59.9
2	female	Netball	173.	63
3	female	Netball	176	66.3
4	female	Netball	170.	60.7
5	female	Netball	183	72.9
6	female	Netball	178.	67.9
7	female	Netball	177.	67.5
8	female	Netball	174.	74.1
9	female	Netball	174.	68.2
10	female	Netball	174.	68.8

```
# i 192 more rows
```

Select-helpers

Other ways to select columns: those whose name:

- ▶ `starts_with` something
- ▶ `ends_with` something
- ▶ `contains` something
- ▶ `matches` a “regular expression”
- ▶ `everything()` select all the columns

Columns whose names begin with S

```
athletes %>% select(starts_with("S"))
```

```
# A tibble: 202 x 3
```

	Sex	Sport	SSF
	<chr>	<chr>	<dbl>
1	female	Netball	49
2	female	Netball	110.
3	female	Netball	89
4	female	Netball	98.3
5	female	Netball	122.
6	female	Netball	90.4
7	female	Netball	107.
8	female	Netball	157.
9	female	Netball	101.
10	female	Netball	126.

```
# i 192 more rows
```

Columns whose names end with C

either uppercase or lowercase:

```
athletes %>% select(ends_with("c"))
```

```
# A tibble: 202 x 3
      RCC    WCC    Hc
  <dbl> <dbl> <dbl>
1  4.56  13.3  42.2
2  4.15    6   38
3  4.16   7.6  37.5
4  4.32   6.4  37.7
5  4.06   5.8  38.7
6  4.12   6.1  36.6
7  4.17    5   37.4
8  3.8    6.6  36.5
9  3.96   5.5  36.3
10 4.44   9.7  41.4
# i 192 more rows
```

Case-sensitive

This works with any of the select-helpers:

```
athletes %>% select(ends_with("C", ignore.case=FALSE))
```

```
# A tibble: 202 x 2
```

```
      RCC    WCC
```

```
  <dbl> <dbl>
```

```
1  4.56  13.3
```

```
2  4.15    6
```

```
3  4.16   7.6
```

```
4  4.32   6.4
```

```
5  4.06   5.8
```

```
6  4.12   6.1
```

```
7  4.17    5
```

```
8  3.8    6.6
```

```
9  3.96   5.5
```

```
10 4.44   9.7
```

```
# i 192 more rows
```

Column names containing letter R

```
athletes %>% select(contains("r"))
```

```
# A tibble: 202 x 3
```

	Sport	RCC	Ferr
	<chr>	<dbl>	<dbl>
1	Netball	4.56	20
2	Netball	4.15	59
3	Netball	4.16	22
4	Netball	4.32	30
5	Netball	4.06	78
6	Netball	4.12	21
7	Netball	4.17	109
8	Netball	3.8	102
9	Netball	3.96	71
10	Netball	4.44	64

```
# i 192 more rows
```

Exactly two characters, ending with T

In regular expression terms, this is `^.t$`:

- ▶ `^` means “start of text”
- ▶ `.` means “exactly one character, but could be anything”
- ▶ `$` means “end of text”.

```
athletes %>% select(matches("^.t$"))
```

```
# A tibble: 202 x 2
```

	Ht	Wt
	<dbl>	<dbl>
1	177.	59.9
2	173.	63
3	176	66.3
4	170.	60.7
5	183	72.9
6	178.	67.9
7	177.	67.5
8	174.	74.1
9	174	68.2

Choosing columns by property

- ▶ Use `where` as with summarizing several columns
- ▶ eg, to choose text columns:

```
athletes %>% select(where(is.character))
```

```
# A tibble: 202 x 2
```

```
  Sex      Sport
```

```
  <chr>   <chr>
```

```
1 female Netball
```

```
2 female Netball
```

```
3 female Netball
```

```
4 female Netball
```

```
5 female Netball
```

```
6 female Netball
```

```
7 female Netball
```

```
8 female Netball
```

```
9 female Netball
```

```
10 female Netball
```

```
# i 192 more rows
```

Choosing rows by number

```
athletes %>% slice(16:25)
```

```
# A tibble: 10 x 13
```

	Sex	Sport	RCC	WCC	Hc	Hg	Ferr	BMI	SSF	``Bfat``	LBM
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	female	Netba~	4.25	10.7	39.5	13.2	127	24.5	157.	26.5	54.5
2	female	Netba~	4.46	10.9	39.7	13.7	102	24.0	116.	23.0	57.2
3	female	Netba~	4.4	9.3	40.4	13.6	86	26.2	182.	30.1	54.4
4	female	Netba~	4.83	8.4	41.8	13.4	40	20.0	71.6	13.9	57.6
5	female	Netba~	4.23	6.9	38.3	12.6	50	25.7	144.	26.6	61.5
6	female	Netba~	4.24	8.4	37.6	12.5	58	25.6	201.	35.5	53.5
7	female	Netba~	3.95	6.6	38.4	12.8	33	19.9	68.9	15.6	54.1
8	female	Netba~	4.03	8.5	37.7	13	51	23.4	104.	19.6	55.4
9	female	BBall	3.96	7.5	37.5	12.3	60	20.6	109.	19.8	63.3
10	female	BBall	4.41	8.3	38.2	12.7	68	20.7	103.	21.3	58.6

```
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Non-consecutive rows

```
athletes %>%  
  slice(10, 13, 17, 42)
```

```
# A tibble: 4 x 13
```

	Sex	Sport	RCC	WCC	Hc	Hg	Ferr	BMI	SSF	``Bfat``	LBM
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	female	Netball	4.44	9.7	41.4	14.1	64	22.8	126.	25.0	51.6
2	female	Netball	4.02	9.1	37.7	12.7	107	23.0	77	18.1	57.3
3	female	Netball	4.46	10.9	39.7	13.7	102	24.0	116.	23.0	57.2
4	female	Row	4.37	8.1	41.8	14.3	53	23.5	98	21.8	63.0

```
# i 2 more variables: Ht <dbl>, Wt <dbl>
```


A random sample of rows

```
athletes %>% slice_sample(n=8)
```

```
# A tibble: 8 x 13
```

	Sex	Sport	RCC	WCC	Hc	Hg	Ferr	BMI	SSF	``Bfat``	LBM
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	male	Field	5.09	8.9	46.3	15.4	44	30.0	71.1	14.0	88
2	male	T400m	5.03	6.6	44.7	15.9	191	19.8	30.9	6.53	59
3	male	Swim	5.13	7.1	46.8	15.9	34	22.5	44.5	8.47	61
4	male	BBall	4.54	5.9	44.4	15.6	97	20.7	41.5	7.16	73
5	female	Swim	4.07	5.9	39.5	13.3	25	20.4	54.6	11.5	52.2
6	male	Tennis	5.03	6.4	42.7	14.3	122	22.0	47.6	8.51	68
7	female	Tennis	4.08	6.6	37.8	12.1	182	20.5	68.3	15.3	51.8
8	female	Gym	4.19	9	39	13.4	69	18.9	43.5	10.2	43.0

```
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Rows for which something is true

```
athletes %>% filter(Sport == "Tennis")
```

```
# A tibble: 11 x 13
```

	Sex	Sport	RCC	WCC	Hc	Hg	Ferr	BMI	SSF	%Bfat`	LBM
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	female	Tennis	4	4.2	36.6	12	57	25.4	109	20.9	56.6
2	female	Tennis	4.4	4	40.8	13.9	73	22.1	98.1	19.6	56.0
3	female	Tennis	4.38	7.9	39.8	13.5	88	21.2	80.6	17.1	46.5
4	female	Tennis	4.08	6.6	37.8	12.1	182	20.5	68.3	15.3	51.8
5	female	Tennis	4.98	6.4	44.8	14.8	80	17.1	47.6	11.1	42.2
6	female	Tennis	5.16	7.2	44.3	14.5	88	18.3	61.9	12.9	48.8
7	female	Tennis	4.66	6.4	40.9	13.9	109	18.4	38.2	8.45	41.9
8	male	Tennis	5.66	8.3	50.2	17.7	38	23.8	56.5	10.0	72
9	male	Tennis	5.03	6.4	42.7	14.3	122	22.0	47.6	8.51	68
10	male	Tennis	4.97	8.8	43	14.9	233	22.3	60.4	11.5	63
11	male	Tennis	5.38	6.3	46	15.7	32	21.1	34.9	6.26	72

```
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

More complicated selections

```
athletes %>% filter(Sport == "Tennis", RCC < 5)
```

```
# A tibble: 7 x 13
```

	Sex	Sport	RCC	WCC	Hc	Hg	Ferr	BMI	SSF
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	female	Tennis	4	4.2	36.6	12	57	25.4	109
2	female	Tennis	4.4	4	40.8	13.9	73	22.1	98.1
3	female	Tennis	4.38	7.9	39.8	13.5	88	21.2	80.6
4	female	Tennis	4.08	6.6	37.8	12.1	182	20.5	68.3
5	female	Tennis	4.98	6.4	44.8	14.8	80	17.1	47.6
6	female	Tennis	4.66	6.4	40.9	13.9	109	18.4	38.2
7	male	Tennis	4.97	8.8	43	14.9	233	22.3	60.4

```
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Another way to do “and”

```
athletes %>% filter(Sport == "Tennis") %>%  
  filter(RCC < 5)
```

```
# A tibble: 7 x 13
```

	Sex	Sport	RCC	WCC	Hc	Hg	Ferr	BMI	SSF
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	female	Tennis	4	4.2	36.6	12	57	25.4	109
2	female	Tennis	4.4	4	40.8	13.9	73	22.1	98.1
3	female	Tennis	4.38	7.9	39.8	13.5	88	21.2	80.6
4	female	Tennis	4.08	6.6	37.8	12.1	182	20.5	68.3
5	female	Tennis	4.98	6.4	44.8	14.8	80	17.1	47.6
6	female	Tennis	4.66	6.4	40.9	13.9	109	18.4	38.2
7	male	Tennis	4.97	8.8	43	14.9	233	22.3	60.4

```
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Either/Or

```
athletes %>% filter(Sport == "Tennis" | RCC > 5)
```

```
# A tibble: 66 x 13
```

	Sex	Sport	RCC	WCC	Hc	Hg	Ferr	BMI	SSF
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	female	Row	5.02	6.4	44.8	15.2	48	19.8	91
2	female	T400m	5.31	9.5	47.1	15.9	29	21.4	57.9
3	female	Field	5.33	9.3	47	15	62	25.3	103.
4	female	TSprnt	5.16	8.2	45.3	14.7	34	20.3	46.1
5	female	Tennis	4	4.2	36.6	12	57	25.4	109
6	female	Tennis	4.4	4	40.8	13.9	73	22.1	98.1
7	female	Tennis	4.38	7.9	39.8	13.5	88	21.2	80.6
8	female	Tennis	4.08	6.6	37.8	12.1	182	20.5	68.3
9	female	Tennis	4.98	6.4	44.8	14.8	80	17.1	47.6
10	female	Tennis	5.16	7.2	44.3	14.5	88	18.3	61.9

```
# i 56 more rows
```

```
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Sorting into order

```
athletes %>% arrange(RCC)
```

```
# A tibble: 202 x 13
```

	Sex	Sport	RCC	WCC	Hc	Hg	Ferr	BMI	SSF
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	female	Netba~	3.8	6.6	36.5	12.4	102	24.4	157.
2	female	Netba~	3.9	6.3	35.9	12.1	78	20.1	70
3	female	T400m	3.9	6	38.9	13.5	16	19.4	48.4
4	female	Row	3.91	7.3	37.6	12.9	43	22.3	126.
5	female	Netba~	3.95	6.6	38.4	12.8	33	19.9	68.9
6	female	Row	3.95	3.3	36.9	12.5	40	24.5	74.9
7	female	Netba~	3.96	5.5	36.3	12.4	71	22.6	101.
8	female	BBall	3.96	7.5	37.5	12.3	60	20.6	109.
9	female	Tennis	4	4.2	36.6	12	57	25.4	109
10	female	Netba~	4.02	9.1	37.7	12.7	107	23.0	77

```
# i 192 more rows  
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Breaking ties by another variable

```
athletes %>% arrange(RCC, BMI)
```

```
# A tibble: 202 x 13
```

	Sex	Sport	RCC	WCC	Hc	Hg	Ferr	BMI	SSF
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	female	Netba~	3.8	6.6	36.5	12.4	102	24.4	157.
2	female	T400m	3.9	6	38.9	13.5	16	19.4	48.4
3	female	Netba~	3.9	6.3	35.9	12.1	78	20.1	70
4	female	Row	3.91	7.3	37.6	12.9	43	22.3	126.
5	female	Netba~	3.95	6.6	38.4	12.8	33	19.9	68.9
6	female	Row	3.95	3.3	36.9	12.5	40	24.5	74.9
7	female	BBall	3.96	7.5	37.5	12.3	60	20.6	109.
8	female	Netba~	3.96	5.5	36.3	12.4	71	22.6	101.
9	female	Tennis	4	4.2	36.6	12	57	25.4	109
10	female	Netba~	4.02	9.1	37.7	12.7	107	23.0	77

```
# i 192 more rows  
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Descending order

```
athletes %>% arrange(desc(BMI))
```

```
# A tibble: 202 x 13
```

	Sex	Sport	RCC	WCC	Hc	Hg	Ferr	BMI	SSF
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	male	Field	5.48	6.2	48.2	16.3	94	34.4	82.7
2	male	Field	4.96	8.3	45.3	15.7	141	33.7	114.
3	male	Field	5.48	4.6	49.4	18	132	32.5	55.7
4	female	Field	4.75	7.5	43.8	15.2	90	31.9	132.
5	male	Field	5.01	8.9	46	15.9	212	30.2	112.
6	male	Field	5.01	8.9	46	15.9	212	30.2	96.9
7	male	Field	5.09	8.9	46.3	15.4	44	30.0	71.1
8	female	Field	4.58	5.8	42.1	14.7	164	28.6	110.
9	female	Field	4.51	9	39.7	14.3	36	28.1	136.
10	male	WPolo	5.34	6.2	49.8	17.2	143	27.8	75.7

```
# i 192 more rows
```

```
# i 2 more variables: Ht <dbl>, Wt <dbl>
```


“The top ones”

```
athletes %>%  
  arrange(desc(Wt)) %>%  
  slice(1:7) %>%  
  select(Sport, Wt)
```

```
# A tibble: 7 x 2
```

```
  Sport      Wt  
  <chr> <dbl>
```

```
1 Field   123.  
2 BBall   114.  
3 Field   111.  
4 Field   108.  
5 Field   103.  
6 WPolo    101  
7 BBall    100.
```

Another way

```
athletes %>%  
  slice_max(order_by = Wt, n=7) %>%  
  select(Sport, Wt)
```

```
# A tibble: 7 x 2
```

```
  Sport      Wt  
  <chr> <dbl>
```

```
1 Field   123.
```

```
2 BBall   114.
```

```
3 Field   111.
```

```
4 Field   108.
```

```
5 Field   103.
```

```
6 WPolo    101
```

```
7 BBall    100.
```

Create new variables from old ones

```
athletes %>%  
  mutate(wt_lb = Wt * 2.2) %>%  
  select(Sport, Sex, Wt, wt_lb) %>%  
  arrange(Wt)
```

```
# A tibble: 202 x 4
```

	Sport	Sex	Wt	wt_lb
	<chr>	<chr>	<dbl>	<dbl>
1	Gym	female	37.8	83.2
2	Gym	female	43.8	96.4
3	Gym	female	45.1	99.2
4	Tennis	female	45.8	101.
5	Tennis	female	47.4	104.
6	Gym	female	47.8	105.
7	T400m	female	49.2	108.
8	Row	female	49.8	110.
9	T400m	female	50.9	112.
10	Netball	female	51.9	114.

```
# ... 100 more rows ...
```

Turning the result into a number

Output is always data frame unless you explicitly turn it into something else, eg. the weight of the heaviest athlete, as a number:

```
athletes %>% arrange(desc(Wt)) %>% pluck("Wt", 1)
```

```
[1] 123.2
```

Or the 20 heaviest weights in descending order:

```
athletes %>%  
  arrange(desc(Wt)) %>%  
  slice(1:20) %>%  
  pluck("Wt")
```

```
[1] 123.20 113.70 111.30 108.20 102.70 101.00 100.20 98.0  
[10] 97.90 97.00 96.90 96.30 94.80 94.80 94.70 94.7  
[19] 94.25 94.20
```

Another way to do the last one

```
athletes %>%  
  arrange(desc(Wt)) %>%  
  slice(1:20) %>%  
  pull("Wt")
```

```
[1] 123.20 113.70 111.30 108.20 102.70 101.00 100.20 98.0  
[10] 97.90 97.00 96.90 96.30 94.80 94.80 94.70 94.7  
[19] 94.25 94.20
```

`pull` grabs the column you name *as a vector* (of whatever it contains).

To find the mean height of the women athletes

Two ways:

```
athletes %>% group_by(Sex) %>% summarize(m = mean(Ht))
```

```
# A tibble: 2 x 2
```

```
  Sex      m  
  <chr> <dbl>
```

```
1 female 175.
```

```
2 male   186.
```

```
athletes %>%
```

```
  filter(Sex == "female") %>%  
  summarize(m = mean(Ht))
```

```
# A tibble: 1 x 1
```

```
      m  
  <dbl>
```

```
1 175.
```

Summary of data selection/arrangement “verbs”

Verb	Purpose
<code>select</code>	Choose columns
<code>slice</code>	Choose rows by number
<code>slice_sample</code>	Choose random rows
<code>slice_max</code>	Choose rows with largest values on a variable (also <code>slice_min</code>)
<code>filter</code>	Choose rows satisfying conditions
<code>arrange</code>	Sort in order by column(s)
<code>mutate</code>	Create new variables
<code>group_by</code>	Create groups to work with
<code>summarize</code>	Calculate summary statistics (by groups if defined)
<code>pluck</code>	Extract items from data frame
<code>pull</code>	Extract a single column from a data frame as a vector

Looking things up in another data frame

- Suppose you are working in the nails department of a hardware store and you find that you have sold these items:

```
my_url <- "http://ritsokiguess.site/datafiles/nail_sales.csv"
sales <- read_csv(my_url)
sales
```

```
# A tibble: 6 x 2
  product_code sales
  <chr>         <dbl>
1 061-5344-6     10
2 161-0090-0      6
3 061-5388-2      2
4 161-0199-4      8
5 061-5375-2      5
6 061-4525-2      3
```


Product descriptions and prices

- ▶ but you don't remember what these product codes are, and you would like to know the total revenue from these sales.
- ▶ Fortunately you found a list of product descriptions and prices:

```
my_url <- "http://ritsokiguess.site/datafiles/nail_desc.csv"
desc <- read_csv(my_url)
desc
```

```
# A tibble: 7 x 5
```

	product_code	description	size	qty	price
	<chr>	<chr>	<chr>	<dbl>	<dbl>
1	061-4525-2	spike nail	"10\""	1	1.49
2	061-5329-4	masonry nail	"1.5\""	112	8.19
3	061-5344-6	finishing nail	"1\""	1298	6.99
4	061-5375-2	roofing nail	"1.25\""	192	6.99
5	061-5388-2	framing nail	"4\""	25	8.19
6	161-0090-0	wood nail	"1\""	25	2.39
7	161-0199-4	panel nail	"1-5/8\""	20	4.69

- ▶ the size values are measured in inches (symbol "), but R uses the same symbol for the start and end of text, so the " representing

The lookup

- ▶ How do you “look up” the product codes to find the product descriptions and prices?
- ▶ `left_join`.

```
sales %>% left_join(desc)
```

```
# A tibble: 6 x 6
```

	product_code	sales	description	size	qty	price
	<chr>	<dbl>	<chr>	<chr>	<dbl>	<dbl>
1	061-5344-6	10	finishing nail	"1\""	1298	6.99
2	161-0090-0	6	wood nail	"1\""	25	2.39
3	061-5388-2	2	framing nail	"4\""	25	8.19
4	161-0199-4	8	panel nail	"1-5/8\""	20	4.69
5	061-5375-2	5	roofing nail	"1.25\""	192	6.99
6	061-4525-2	3	spike nail	"10\""	1	1.49

What we have

- ▶ this looks up all the rows in the *first* dataframe that are also in the *second*.
- ▶ by default matches all columns with same name in two dataframes (product_code here)
- ▶ get *all* columns in *both* dataframes. The rows are the ones for that product_code.

So now can work out how much the total revenue was:

```
sales %>% left_join(desc) %>%  
  mutate(product_revenue = sales*price) %>%  
  summarize(total_revenue = sum(product_revenue))
```

```
# A tibble: 1 x 1  
  total_revenue  
      <dbl>  
1         178.
```

More comments

- ▶ if any product codes are not matched, you get NA in the added columns
- ▶ anything in the *second* dataframe that was not in the first does not appear (here, any products that were not sold)
- ▶ other variations (examples follow):
 - ▶ if there are two columns with the same name in the two dataframes, and you only want to match on one, use `by` with one column name
 - ▶ if the columns you want to look up have different names in the two dataframes, use `by` with a “named list”

Matching on only some matching names

- ▶ Suppose the sales dataframe *also* had a column qty (which was the quantity sold):

```
sales %>% rename("qty"="sales") -> sales1
sales1
```

```
# A tibble: 6 x 2
  product_code qty
  <chr>      <dbl>
1 061-5344-6    10
2 161-0090-0     6
3 061-5388-2     2
4 161-0199-4     8
5 061-5375-2     5
6 061-4525-2     3
```

- ▶ The qty in sales1 is the quantity sold, but the qty in desc is the number of nails in a package. These should *not* be matched: they are different things.

Matching only on product code

```
sales1 %>%  
  left_join(desc, by = "product_code")
```

```
# A tibble: 6 x 6
```

	product_code	qty.x	description	size	qty.y	price
	<chr>	<dbl>	<chr>	<chr>	<dbl>	<dbl>
1	061-5344-6	10	finishing nail	"1\""	1298	6.99
2	161-0090-0	6	wood nail	"1\""	25	2.39
3	061-5388-2	2	framing nail	"4\""	25	8.19
4	161-0199-4	8	panel nail	"1-5/8\""	20	4.69
5	061-5375-2	5	roofing nail	"1.25\""	192	6.99
6	061-4525-2	3	spike nail	"10\""	1	1.49

► Get qty.x (from sales1) and qty.y (from desc).

Matching on different names 1/2

- Suppose the product code in sales was just code:

```
sales %>% rename("code" = "product_code") -> sales2  
sales2
```

```
# A tibble: 6 x 2  
  code      sales  
  <chr>    <dbl>  
1 061-5344-6    10  
2 161-0090-0     6  
3 061-5388-2     2  
4 161-0199-4     8  
5 061-5375-2     5  
6 061-4525-2     3
```

- How to match the two product codes that have different names?

Matching on different names 2/2

► Use by, but like this:

```
sales2 %>%  
  left_join(desc, by = c("code"="product_code"))
```

```
# A tibble: 6 x 6
```

	code	sales	description	size	qty	price
	<chr>	<dbl>	<chr>	<chr>	<dbl>	<dbl>
1	061-5344-6	10	finishing nail	"1\""	1298	6.99
2	161-0090-0	6	wood nail	"1\""	25	2.39
3	061-5388-2	2	framing nail	"4\""	25	8.19
4	161-0199-4	8	panel nail	"1-5/8\""	20	4.69
5	061-5375-2	5	roofing nail	"1.25\""	192	6.99
6	061-4525-2	3	spike nail	"10\""	1	1.49

Other types of join

- ▶ `right_join`: interchanges roles, looking up keys from second dataframe in first.
- ▶ `anti_join`: give me all the rows in the first dataframe that are *not* in the second. (Use this eg. to see whether the product descriptions are incomplete.)
- ▶ `full_join`: give me all the rows in both dataframes, with missings as needed.

Full join here

```
sales %>% full_join(desc)
```

```
# A tibble: 7 x 6
```

	product_code <chr>	sales <dbl>	description <chr>	size <chr>	qty <dbl>	price <dbl>
1	061-5344-6	10	finishing nail	"1\""	1298	6.99
2	161-0090-0	6	wood nail	"1\""	25	2.39
3	061-5388-2	2	framing nail	"4\""	25	8.19
4	161-0199-4	8	panel nail	"1-5/8\""	20	4.69
5	061-5375-2	5	roofing nail	"1.25\""	192	6.99
6	061-4525-2	3	spike nail	"10\""	1	1.49
7	061-5329-4	NA	masonry nail	"1.5\""	112	8.19

- The missing sales for “masonry nail” says that it was in the lookup table desc, but we didn’t sell any.

The same thing, but with anti_join

Anything in first df but not in second?

```
desc %>% anti_join(sales)
```

```
# A tibble: 1 x 5
```

	product_code	description	size	qty	price
	<chr>	<chr>	<chr>	<dbl>	<dbl>
1	061-5329-4	masonry nail	"1.5\\""	112	8.19

Masonry nails are the only thing in our product description file that we did not sell any of.