Choosing things in dataframes

Packages

The usual:

```
library(tidyverse)
```

Doing things with data frames

Let's go back to our Australian athletes:
```
athletes
```

```
# A tibble: 202 x 13
   Sex    Sport   RCC   WCC    Hc    Hg  Ferr   BMI   SSF `%Bfat`   LBM
   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl> <dbl>
 1 female Netba~ 4.56 13.3   42.2  13.6    20  19.2  49      11.3  53.1
 2 female Netba~ 4.15  6     38    12.7    59  21.2 110.     25.3  47.1
 3 female Netba~ 4.16  7.6   37.5  12.3    22  21.4  89      19.4  53.4
 4 female Netba~ 4.32  6.4   37.7  12.3    30  21.0  98.3    19.6  48.8
 5 female Netba~ 4.06  5.8   38.7  12.8    78  21.8 122.     23.1  56.0
 6 female Netba~ 4.12  6.1   36.6  11.8    21  21.4  90.4    16.9  56.4
 7 female Netba~ 4.17  5     37.4  12.7   109  21.5 107.     21.3  53.1
 8 female Netba~ 3.8   6.6   36.5  12.4   102  24.4 157.     26.6  54.4
 9 female Netba~ 3.96  5.5   36.3  12.4    71  22.6 101.     17.9  56.0
10 female Netba~ 4.44  9.7   41.4  14.1    64  22.8 126.     25.0  51.6
# i 192 more rows
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Choosing a column

```
athletes %>% select(Sport)

# A tibble: 202 x 1
   Sport
   <chr>
 1 Netball
 2 Netball
 3 Netball
 4 Netball
 5 Netball
 6 Netball
 7 Netball
 8 Netball
 9 Netball
10 Netball
# i 192 more rows
```

Choosing several columns

```
athletes %>% select(Sport, Hg, BMI)

# A tibble: 202 x 3
   Sport      Hg   BMI
   <chr>   <dbl> <dbl>
 1 Netball  13.6  19.2
 2 Netball  12.7  21.2
 3 Netball  12.3  21.4
 4 Netball  12.3  21.0
 5 Netball  12.8  21.8
 6 Netball  11.8  21.4
 7 Netball  12.7  21.5
 8 Netball  12.4  24.4
 9 Netball  12.4  22.6
10 Netball  14.1  22.8
# i 192 more rows
```

Choosing consecutive columns

```
athletes %>% select(Sex:WCC)
```

```
# A tibble: 202 x 4
   Sex    Sport     RCC   WCC
   <chr>  <chr>   <dbl> <dbl>
 1 female Netball  4.56  13.3
 2 female Netball  4.15   6
 3 female Netball  4.16   7.6
 4 female Netball  4.32   6.4
 5 female Netball  4.06   5.8
 6 female Netball  4.12   6.1
 7 female Netball  4.17   5
 8 female Netball  3.8    6.6
 9 female Netball  3.96   5.5
10 female Netball  4.44   9.7
# i 192 more rows
```

Choosing all-but some columns

```
athletes %>% select(-(RCC:LBM))
```

```
# A tibble: 202 x 4
   Sex    Sport     Ht    Wt
   <chr>  <chr>  <dbl> <dbl>
 1 female Netball  177.  59.9
 2 female Netball  173.  63
 3 female Netball  176   66.3
 4 female Netball  170.  60.7
 5 female Netball  183   72.9
 6 female Netball  178.  67.9
 7 female Netball  177.  67.5
 8 female Netball  174.  74.1
 9 female Netball  174.  68.2
10 female Netball  174.  68.8
# i 192 more rows
```

Select-helpers

Other ways to select columns: those whose name:

▶ `starts_with` something
▶ `ends_with` something
▶ `contains` something
▶ `matches` a "regular expression"
▶ `everything()` select all the columns

Columns whose names begin with S

```
athletes %>% select(starts_with("S"))
```

```
# A tibble: 202 x 3
   Sex    Sport      SSF
   <chr>  <chr>    <dbl>
 1 female Netball     49
 2 female Netball    110.
 3 female Netball     89
 4 female Netball     98.3
 5 female Netball    122.
 6 female Netball     90.4
 7 female Netball    107.
 8 female Netball    157.
 9 female Netball    101.
10 female Netball    126.
# i 192 more rows
```

Columns whose names end with C

either uppercase or lowercase:

```
athletes %>% select(ends_with("c"))
```

```
# A tibble: 202 x 3
      RCC   WCC    Hc
    <dbl> <dbl> <dbl>
 1  4.56  13.3  42.2
 2  4.15   6    38
 3  4.16   7.6  37.5
 4  4.32   6.4  37.7
 5  4.06   5.8  38.7
 6  4.12   6.1  36.6
 7  4.17   5    37.4
 8  3.8    6.6  36.5
 9  3.96   5.5  36.3
10  4.44   9.7  41.4
# i 192 more rows
```

Case-sensitive

This works with any of the select-helpers:

```
athletes %>% select(ends_with("C", ignore.case=FALSE))
```

```
# A tibble: 202 x 2
     RCC   WCC
   <dbl> <dbl>
 1  4.56  13.3
 2  4.15   6
 3  4.16   7.6
 4  4.32   6.4
 5  4.06   5.8
 6  4.12   6.1
 7  4.17   5
 8  3.8    6.6
 9  3.96   5.5
10  4.44   9.7
# i 192 more rows
```

Column names containing letter R

```
athletes %>% select(contains("r"))
```

```
# A tibble: 202 x 3
   Sport     RCC   Ferr
   <chr>    <dbl>  <dbl>
 1 Netball   4.56    20
 2 Netball   4.15    59
 3 Netball   4.16    22
 4 Netball   4.32    30
 5 Netball   4.06    78
 6 Netball   4.12    21
 7 Netball   4.17   109
 8 Netball   3.8    102
 9 Netball   3.96    71
10 Netball   4.44    64
# i 192 more rows
```

Exactly two characters, ending with T

In regular expression terms, this is `^.t$`:

▶ `^` means "start of text"
▶ `.` means "exactly one character, but could be anything"
▶ `$` means "end of text".

```
athletes %>% select(matches("^.t$"))
```

```
# A tibble: 202 x 2
      Ht    Wt
   <dbl> <dbl>
 1  177.  59.9
 2  173.  63
 3  176   66.3
 4  170.  60.7
 5  183   72.9
 6  178.  67.9
 7  177.  67.5
 8  174.  74.1
 9  174.  68.2
```

Choosing columns by property

▶ Use where as with summarizing several columns

▶ eg, to choose text columns:

```
athletes %>% select(where(is.character))
```

```
# A tibble: 202 x 2
   Sex    Sport
   <chr>  <chr>
 1 female Netball
 2 female Netball
 3 female Netball
 4 female Netball
 5 female Netball
 6 female Netball
 7 female Netball
 8 female Netball
 9 female Netball
10 female Netball
# i 192 more rows
```

Choosing rows by number

```
athletes %>% slice(16:25)
```

```
# A tibble: 10 x 13
   Sex    Sport   RCC   WCC    Hc    Hg  Ferr   BMI   SSF `%Bfat`   LBM
   <chr>  <chr>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl> <dbl>
 1 female Netba~  4.25  10.7  39.5  13.2   127  24.5 157.     26.5  54.5
 2 female Netba~  4.46  10.9  39.7  13.7   102  24.0 116.     23.0  57.2
 3 female Netba~  4.4    9.3  40.4  13.6    86  26.2 182.     30.1  54.4
 4 female Netba~  4.83   8.4  41.8  13.4    40  20.0  71.6    13.9  57.6
 5 female Netba~  4.23   6.9  38.3  12.6    50  25.7 144.     26.6  61.5
 6 female Netba~  4.24   8.4  37.6  12.5    58  25.6 201.     35.5  53.5
 7 female Netba~  3.95   6.6  38.4  12.8    33  19.9  68.9    15.6  54.1
 8 female Netba~  4.03   8.5  37.7  13      51  23.4 104.     19.6  55.4
 9 female BBall   3.96   7.5  37.5  12.3    60  20.6 109.     19.8  63.3
10 female BBall   4.41   8.3  38.2  12.7    68  20.7 103.     21.3  58.6
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Non-consecutive rows

```
athletes %>%
  slice(10, 13, 17, 42)
```

```
# A tibble: 4 x 13
  Sex    Sport     RCC   WCC    Hc    Hg  Ferr   BMI   SSF `%Bfat`   LBM
  <chr>  <chr>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl> <dbl>
1 female Netball  4.44   9.7  41.4  14.1    64  22.8  126.    25.0  51.6
2 female Netball  4.02   9.1  37.7  12.7   107  23.0   77     18.1  57.3
3 female Netball  4.46  10.9  39.7  13.7   102  24.0  116.    23.0  57.2
4 female Row      4.37   8.1  41.8  14.3    53  23.5   98     21.8  63.0
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

A random sample of rows

```
athletes %>% slice_sample(n=8)
```

```
# A tibble: 8 x 13
  Sex    Sport     RCC   WCC    Hc    Hg  Ferr   BMI   SSF `%Bfat`   LBM
  <chr>  <chr>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl> <dbl>
1 female Gym      4.19   9    39    13.4    69  18.9  43.5    10.2  43.0
2 female TSprnt   5.16   8.2  45.3  14.7    34  20.3  46.1    10.2  51.5
3 female Tennis   4.98   6.4  44.8  14.8    80  17.1  47.6    11.1  42.2
4 female Netball  4.39   9.6  38.3  12.5    39  24.6 149.     26.8  54.2
5 male   Field    5.11   9.6  48.2  16.7   103  27.4  65.9    11.7  83
6 female Row      4.49   7.6  41.8  14.4    71  24.2 111.     24.7  56.5
7 male   Row      4.87   8.2  43.8  15     130  23.6  49.2     9    78
8 female Field    4.51   9    39.7  14.3    36  28.1 136.     24.9  63.0
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Rows for which something is true

```
athletes %>% filter(Sport == "Tennis")
```

```
# A tibble: 11 x 13
   Sex    Sport   RCC   WCC    Hc    Hg  Ferr   BMI   SSF `%Bfat`   LBM
   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl> <dbl>
 1 female Tennis  4     4.2  36.6  12      57  25.4 109     20.9   56.6
 2 female Tennis  4.4   4    40.8  13.9    73  22.1  98.1   19.6   56.0
 3 female Tennis  4.38  7.9  39.8  13.5    88  21.2  80.6   17.1   46.5
 4 female Tennis  4.08  6.6  37.8  12.1   182  20.5  68.3   15.3   51.8
 5 female Tennis  4.98  6.4  44.8  14.8    80  17.1  47.6   11.1   42.2
 6 female Tennis  5.16  7.2  44.3  14.5    88  18.3  61.9   12.9   48.8
 7 female Tennis  4.66  6.4  40.9  13.9   109  18.4  38.2    8.45  41.9
 8 male   Tennis  5.66  8.3  50.2  17.7    38  23.8  56.5   10.0   72
 9 male   Tennis  5.03  6.4  42.7  14.3   122  22.0  47.6    8.51  68
10 male   Tennis  4.97  8.8  43    14.9   233  22.3  60.4   11.5   63
11 male   Tennis  5.38  6.3  46    15.7    32  21.1  34.9    6.26  72
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

More complicated selections

```
athletes %>% filter(Sport == "Tennis", RCC < 5)

# A tibble: 7 x 13
  Sex    Sport   RCC   WCC    Hc    Hg  Ferr   BMI   SSF
  <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 female Tennis  4     4.2  36.6  12     57  25.4 109
2 female Tennis  4.4   4    40.8  13.9   73  22.1  98.1
3 female Tennis  4.38  7.9  39.8  13.5   88  21.2  80.6
4 female Tennis  4.08  6.6  37.8  12.1  182  20.5  68.3
5 female Tennis  4.98  6.4  44.8  14.8   80  17.1  47.6
6 female Tennis  4.66  6.4  40.9  13.9  109  18.4  38.2
7 male   Tennis  4.97  8.8  43    14.9  233  22.3  60.4
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Another way to do "and"

```
athletes %>% filter(Sport == "Tennis") %>%
  filter(RCC < 5)

# A tibble: 7 x 13
  Sex    Sport  RCC   WCC   Hc    Hg   Ferr   BMI   SSF
  <chr>  <chr>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 female Tennis 4     4.2   36.6  12     57   25.4 109
2 female Tennis 4.4   4     40.8  13.9   73   22.1  98.1
3 female Tennis 4.38  7.9   39.8  13.5   88   21.2  80.6
4 female Tennis 4.08  6.6   37.8  12.1  182   20.5  68.3
5 female Tennis 4.98  6.4   44.8  14.8   80   17.1  47.6
6 female Tennis 4.66  6.4   40.9  13.9  109   18.4  38.2
7 male   Tennis 4.97  8.8   43    14.9  233   22.3  60.4
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Either/Or

```r
athletes %>% filter(Sport == "Tennis" | RCC > 5)
```

```
# A tibble: 66 x 13
   Sex    Sport   RCC   WCC    Hc    Hg  Ferr   BMI   SSF
   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
 1 female Row    5.02   6.4  44.8  15.2    48  19.8  91
 2 female T400m  5.31   9.5  47.1  15.9    29  21.4  57.9
 3 female Field  5.33   9.3  47    15      62  25.3 103.
 4 female TSprnt 5.16   8.2  45.3  14.7    34  20.3  46.1
 5 female Tennis 4      4.2  36.6  12      57  25.4 109
 6 female Tennis 4.4    4    40.8  13.9    73  22.1  98.1
 7 female Tennis 4.38   7.9  39.8  13.5    88  21.2  80.6
 8 female Tennis 4.08   6.6  37.8  12.1   182  20.5  68.3
 9 female Tennis 4.98   6.4  44.8  14.8    80  17.1  47.6
10 female Tennis 5.16   7.2  44.3  14.5    88  18.3  61.9
# i 56 more rows
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Sorting into order

```
athletes %>% arrange(RCC)
```

```
# A tibble: 202 x 13
   Sex    Sport   RCC   WCC    Hc    Hg  Ferr   BMI   SSF
   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
 1 female Netba~  3.8   6.6  36.5  12.4   102  24.4 157.
 2 female Netba~  3.9   6.3  35.9  12.1    78  20.1  70
 3 female T400m   3.9   6    38.9  13.5    16  19.4  48.4
 4 female Row     3.91  7.3  37.6  12.9    43  22.3 126.
 5 female Netba~  3.95  6.6  38.4  12.8    33  19.9  68.9
 6 female Row     3.95  3.3  36.9  12.5    40  24.5  74.9
 7 female Netba~  3.96  5.5  36.3  12.4    71  22.6 101.
 8 female BBall   3.96  7.5  37.5  12.3    60  20.6 109.
 9 female Tennis  4     4.2  36.6  12      57  25.4 109
10 female Netba~  4.02  9.1  37.7  12.7   107  23.0  77
# i 192 more rows
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Breaking ties by another variable

```
athletes %>% arrange(RCC, BMI)
```

```
# A tibble: 202 x 13
   Sex    Sport  RCC   WCC   Hc    Hg   Ferr   BMI   SSF
   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
 1 female Netba~  3.8   6.6  36.5  12.4   102  24.4  157.
 2 female T400m   3.9   6    38.9  13.5    16  19.4   48.4
 3 female Netba~  3.9   6.3  35.9  12.1    78  20.1   70
 4 female Row     3.91  7.3  37.6  12.9    43  22.3  126.
 5 female Netba~  3.95  6.6  38.4  12.8    33  19.9   68.9
 6 female Row     3.95  3.3  36.9  12.5    40  24.5   74.9
 7 female BBall   3.96  7.5  37.5  12.3    60  20.6  109.
 8 female Netba~  3.96  5.5  36.3  12.4    71  22.6  101.
 9 female Tennis  4     4.2  36.6  12      57  25.4  109
10 female Netba~  4.02  9.1  37.7  12.7   107  23.0   77
# i 192 more rows
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

Descending order

```
athletes %>% arrange(desc(BMI))
```

```
# A tibble: 202 x 13
   Sex    Sport   RCC   WCC    Hc    Hg  Ferr   BMI   SSF
   <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
 1 male   Field  5.48   6.2  48.2  16.3    94  34.4  82.7
 2 male   Field  4.96   8.3  45.3  15.7   141  33.7 114.
 3 male   Field  5.48   4.6  49.4  18     132  32.5  55.7
 4 female Field  4.75   7.5  43.8  15.2    90  31.9 132.
 5 male   Field  5.01   8.9  46    15.9   212  30.2 112.
 6 male   Field  5.01   8.9  46    15.9   212  30.2  96.9
 7 male   Field  5.09   8.9  46.3  15.4    44  30.0  71.1
 8 female Field  4.58   5.8  42.1  14.7   164  28.6 110.
 9 female Field  4.51   9    39.7  14.3    36  28.1 136.
10 male   WPolo  5.34   6.2  49.8  17.2   143  27.8  75.7
# i 192 more rows
# i 2 more variables: Ht <dbl>, Wt <dbl>
```

"The top ones"

```
athletes %>%
  arrange(desc(Wt)) %>%
  slice(1:7) %>%
  select(Sport, Wt)
```

```
# A tibble: 7 x 2
  Sport    Wt
  <chr> <dbl>
1 Field  123.
2 BBall  114.
3 Field  111.
4 Field  108.
5 Field  103.
6 WPolo  101
7 BBall  100.
```

Another way

```
athletes %>%
  slice_max(order_by = Wt, n=7) %>%
  select(Sport, Wt)

# A tibble: 7 x 2
  Sport    Wt
  <chr> <dbl>
1 Field  123.
2 BBall  114.
3 Field  111.
4 Field  108.
5 Field  103.
6 WPolo  101
7 BBall  100.
```

Create new variables from old ones

```
athletes %>%
  mutate(wt_lb = Wt * 2.2) %>%
  select(Sport, Sex, Wt, wt_lb) %>%
  arrange(Wt)
```

```
# A tibble: 202 x 4
   Sport   Sex       Wt wt_lb
   <chr>   <chr>  <dbl> <dbl>
 1 Gym     female  37.8  83.2
 2 Gym     female  43.8  96.4
 3 Gym     female  45.1  99.2
 4 Tennis  female  45.8 101.
 5 Tennis  female  47.4 104.
 6 Gym     female  47.8 105.
 7 T400m   female  49.2 108.
 8 Row     female  49.8 110.
 9 T400m   female  50.9 112.
10 Netball female  51.9 114.
# i 192 more rows
```

Turning the result into a number

Output is always data frame unless you explicitly turn it into something else, eg. the weight of the heaviest athlete, as a number:

```
athletes %>% arrange(desc(Wt)) %>% pluck("Wt", 1)
```

```
[1] 123.2
```

Or the 20 heaviest weights in descending order:

```
athletes %>%
  arrange(desc(Wt)) %>%
  slice(1:20) %>%
  pluck("Wt")
```

```
 [1] 123.20 113.70 111.30 108.20 102.70 101.00 100.20  98.0
[10]  97.90  97.00  96.90  96.30  94.80  94.80  94.70  94.7
[19]  94.25  94.20
```

Another way to do the last one

```
athletes %>%
  arrange(desc(Wt)) %>%
  slice(1:20) %>%
  pull("Wt")
```

```
 [1] 123.20 113.70 111.30 108.20 102.70 101.00 100.20  98.0
[10]  97.90  97.00  96.90  96.30  94.80  94.80  94.70  94.7
[19]  94.25  94.20
```

pull grabs the column you name *as a vector* (of whatever it contains).

To find the mean height of the women athletes

Two ways:

```
athletes %>% group_by(Sex) %>% summarize(m = mean(Ht))
```

```
# A tibble: 2 x 2
  Sex        m
  <chr>  <dbl>
1 female  175.
2 male    186.
```

```
athletes %>%
  filter(Sex == "female") %>%
  summarize(m = mean(Ht))
```

```
# A tibble: 1 x 1
      m
  <dbl>
1  175.
```

Summary of data selection/arrangement "verbs"

Verb	Purpose
select	Choose columns
slice	Choose rows by number
slice_sample	Choose random rows
slice_max	Choose rows with largest values on a variable (also slice_min)
filter	Choose rows satisfying conditions
arrange	Sort in order by column(s)
mutate	Create new variables
group_by	Create groups to work with
summarize	Calculate summary statistics (by groups if defined)
pluck	Extract items from data frame
pull	Extract a single column from a data frame as a vector

Looking things up in another data frame

▶ Suppose you are working in the nails department of a hardware store and you find that you have sold these items:

```
my_url <- "http://ritsokiguess.site/datafiles/nail_sales.cs
sales <- read_csv(my_url)
sales
```

```
# A tibble: 6 x 2
  product_code sales
  <chr>        <dbl>
1 061-5344-6      10
2 161-0090-0       6
3 061-5388-2       2
4 161-0199-4       8
5 061-5375-2       5
6 061-4525-2       3
```

Product descriptions and prices

▶ but you don't remember what these product codes are, and you would like to know the total revenue from these sales.

▶ Fortunately you found a list of product descriptions and prices:

```
my_url <- "http://ritsokiguess.site/datafiles/nail_desc.csv"
desc <- read_csv(my_url)
desc
```

```
# A tibble: 7 x 5
  product_code description   size      qty  price
  <chr>        <chr>         <chr>    <dbl>  <dbl>
1 061-4525-2   spike nail     "10\""      1   1.49
2 061-5329-4   masonry nail   "1.5\""   112   8.19
3 061-5344-6   finishing nail "1\""    1298   6.99
4 061-5375-2   roofing nail   "1.25\""  192   6.99
5 061-5388-2   framing nail   "4\""      25   8.19
6 161-0090-0   wood nail      "1\""      25   2.39
7 161-0199-4   panel nail     "1-5/8\""  20   4.69
```

▶ the size values are measured in inches (symbol "), but R uses the same symbol for the start and end of text, so the " representing

The lookup

▶ How do you "look up" the product codes to find the product descriptions and prices?

▶ `left_join`.

```
sales %>% left_join(desc)
```

```
# A tibble: 6 x 6
  product_code sales description   size        qty price
  <chr>        <dbl> <chr>         <chr>      <dbl> <dbl>
1 061-5344-6      10 finishing nail "1\""      1298  6.99
2 161-0090-0       6 wood nail      "1\""        25  2.39
3 061-5388-2       2 framing nail   "4\""        25  8.19
4 161-0199-4       8 panel nail     "1-5/8\""    20  4.69
5 061-5375-2       5 roofing nail   "1.25\""    192  6.99
6 061-4525-2       3 spike nail     "10\""        1  1.49
```

What we have

▶ this looks up all the rows in the *first* dataframe that are also in the *second*.

▶ by default matches all columns with same name in two dataframes (`product_code` here)

▶ get *all* columns in *both* dataframes. The rows are the ones for that `product_code`.

So now can work out how much the total revenue was:

```
sales %>% left_join(desc) %>%
  mutate(product_revenue = sales*price) %>%
  summarize(total_revenue = sum(product_revenue))
```

```
# A tibble: 1 x 1
  total_revenue
          <dbl>
1          178.
```

More comments

- if any product codes are not matched, you get NA in the added columns
- anything in the *second* dataframe that was not in the first does not appear (here, any products that were not sold)
- other variations (examples follow):
    - if there are two columns with the same name in the two dataframes, and you only want to match on one, use by with one column name
    - if the columns you want to look up have different names in the two dataframes, use by with a "named list"

Matching on only some matching names

▶ Suppose the sales dataframe *also* had a column qty (which was the quantity sold):

```
sales %>% rename("qty"="sales") -> sales1
sales1
```

```
# A tibble: 6 x 2
  product_code   qty
  <chr>        <dbl>
1 061-5344-6      10
2 161-0090-0       6
3 061-5388-2       2
4 161-0199-4       8
5 061-5375-2       5
6 061-4525-2       3
```

▶ The qty in sales1 is the quantity sold, but the qty in desc is the number of nails in a package. These should *not* be matched: they are different things.

Matching only on product code

```
sales1 %>%
  left_join(desc, by = "product_code")

# A tibble: 6 x 6
  product_code qty.x description    size     qty.y price
  <chr>        <dbl> <chr>          <chr>    <dbl> <dbl>
1 061-5344-6      10 finishing nail "1\""     1298  6.99
2 161-0090-0       6 wood nail      "1\""       25  2.39
3 061-5388-2       2 framing nail   "4\""       25  8.19
4 161-0199-4       8 panel nail     "1-5/8\""   20  4.69
5 061-5375-2       5 roofing nail   "1.25\""   192  6.99
6 061-4525-2       3 spike nail     "10\""       1  1.49
```

▶ Get qty.x (from sales1) and qty.y (from desc).

Matching on different names 1/2

▶ Suppose the product code in `sales` was just code:

```
sales %>% rename("code" = "product_code") -> sales2
sales2
```

```
# A tibble: 6 x 2
  code      sales
  <chr>     <dbl>
1 061-5344-6   10
2 161-0090-0    6
3 061-5388-2    2
4 161-0199-4    8
5 061-5375-2    5
6 061-4525-2    3
```

▶ How to match the two product codes that have different names?

Matching on different names 2/2

▶ Use by, but like this:

```
sales2 %>%
  left_join(desc, by = c("code"="product_code"))
```

```
# A tibble: 6 x 6
  code       sales description    size        qty price
  <chr>      <dbl> <chr>          <chr>      <dbl> <dbl>
1 061-5344-6    10 finishing nail "1\""       1298  6.99
2 161-0090-0     6 wood nail      "1\""         25  2.39
3 061-5388-2     2 framing nail   "4\""         25  8.19
4 161-0199-4     8 panel nail     "1-5/8\""     20  4.69
5 061-5375-2     5 roofing nail   "1.25\""     192  6.99
6 061-4525-2     3 spike nail     "10\""         1  1.49
```

Other types of join

▶ `right_join`: interchanges roles, looking up keys from second dataframe in first.

▶ `anti_join`: give me all the rows in the first dataframe that are *not* in the second. (Use this eg. to see whether the product descriptions are incomplete.)

▶ `full_join`: give me all the rows in both dataframes, with missings as needed.

Full join here

```
sales %>% full_join(desc)

# A tibble: 7 x 6
  product_code sales description    size       qty price
  <chr>        <dbl> <chr>          <chr>    <dbl> <dbl>
1 061-5344-6     10  finishing nail "1\""     1298  6.99
2 161-0090-0      6  wood nail      "1\""       25  2.39
3 061-5388-2      2  framing nail   "4\""       25  8.19
4 161-0199-4      8  panel nail     "1-5/8\""   20  4.69
5 061-5375-2      5  roofing nail   "1.25\""   192  6.99
6 061-4525-2      3  spike nail     "10\""       1  1.49
7 061-5329-4     NA  masonry nail   "1.5\""    112  8.19
```

▶ The missing `sales` for "masonry nail" says that it was in the lookup table `desc`, but we didn't sell any.

The same thing, but with `anti_join`

Anything in first df but not in second?

```
desc %>% anti_join(sales)
```

```
# A tibble: 1 x 5
  product_code description  size      qty price
  <chr>        <chr>        <chr>   <dbl> <dbl>
1 061-5329-4   masonry nail "1.5\""   112  8.19
```

Masonry nails are the only thing in our product description file that
we did not sell any of.