

# Cluster analysis

## Cluster Analysis

- One side-effect of discriminant analysis: could draw picture of data (if 1st 2s LDs told most of story) and see which individuals “close” to each other.
- Discriminant analysis requires knowledge of groups.
- Without knowledge of groups, use *cluster analysis*: see which individuals close together, which groups suggested by data.
- Idea: see how individuals group into “clusters” of nearby individuals.
- Base on “dissimilarities” between individuals.
- Or base on standard deviations and correlations between variables (assesses dissimilarity behind scenes).

## Packages

```
library(MASS) # for lda later
library(tidyverse)
library(spatstat) # for crossdist later
library(ggrepel)
library(conflicted)
conflict_prefer("select", "dplyr")
conflict_prefer("filter", "dplyr")
```

English	Norwegian	Danish	Dutch	German
---------	-----------	--------	-------	--------

## One to ten in 11 languages

	English	Norwegian	Danish	Dutch	German
1	one	en	en	een	eins
2	two	to	to	twee	zwei
3	three	tre	tre	drie	drei
4	four	fire	fire	vier	vier
5	five	fem	fem	vijf	funf
6	six	seks	seks	zes	sechs
7	seven	sju	syv	zeven	sieben
8	eight	atte	otte	acht	acht
9	nine	ni	ni	negen	neun
10	ten	ti	ti	tien	zehn

## One to ten

	French	Spanish	Italian	Polish	Hungarian	Finnish
1	un	uno	uno	jeden	egy	yksi
2	deux	dos	due	dwa	ketto	kaksi
3	trois	tres	tre	trzy	harom	kolme
4	quatre	cuatro	quattro	cztery	negy	nelja
5	cinq	cinco	cinque	piec	ot	viisi
6	six	seis	sei	szesc	hat	kuusi
7	sept	siete	sette	siedem	het	seitseman
8	huit	ocho	otto	osiem	nyolc	kahdeksan
9	neuf	nueve	nove	dziewiec	kilenc	yhdeksan
10	dix	diez	dieci	dziesiec	tiz	kymmenen

## Dissimilarities and languages example


- Can define dissimilarities how you like (whatever makes sense in application).
- Sometimes defining “similarity” makes more sense; can turn this into dissimilarity by subtracting from some maximum.

- Example: numbers 1–10 in various European languages. Define similarity between two languages by counting how often the same number has a name starting with the same letter (and dissimilarity by how often number has names starting with different letter).
- Crude (doesn't even look at most of the words), but see how effective.

## Two kinds of cluster analysis

- Looking at process of forming clusters (of similar languages): **hierarchical cluster analysis** (`hclust`).
- Start with each individual in cluster by itself.
- Join “closest” clusters one by one until all individuals in one cluster.
- How to define closeness of two *clusters*? Not obvious, investigate in a moment.
- Know how many clusters: which division into that many clusters is “best” for individuals? **K-means clustering** (`kmeans`).


## Two made-up clusters



How to measure distance between set of red points and set of blue ones?

### Single-linkage distance


Find the red point and the blue point that are closest together:



Single-linkage distance between 2 clusters is distance between their closest points.

### Complete linkage


Find the red and blue points that are farthest apart:



Complete-linkage distance is distance between farthest points.

### Ward's method


Work out mean of each cluster and join point to its mean:



Work out (i) sum of squared distances of points from means.

## Ward's method part 2

Now imagine combining the two clusters and working out overall mean. Join each point to this mean:



Calc sum of squared distances (ii) of points to combined mean.

## Ward's method part 3

- Sum of squares (ii) will be bigger than (i) (points closer to own cluster mean than combined mean).
- Ward's distance is (ii) minus (i).
- Think of as “cost” of combining clusters:
- if clusters close together, (ii) only a little larger than (i)
- if clusters far apart, (ii) a lot larger than (i) (as in example).

## Hierarchical clustering revisited

- Single linkage, complete linkage, Ward are ways of measuring closeness of clusters.
- Use them, starting with each observation in own cluster, to repeatedly combine two closest clusters until all points in one cluster.
- They will give different answers (clustering stories).

- Single linkage tends to make “stringy” clusters because clusters can be very different apart from two closest points.
- Complete linkage insists on whole clusters being similar.
- Ward tends to form many small clusters first.

## Dissimilarity data in R

Dissimilarities for language data were how many number names had *different* first letter:

```
my_url <- "http://ritsokiguess.site/datafiles/languages.txt"
(number.d <- read_table(my_url))
```

```
# A tibble: 11 x 12
  la      en    no    dk    nl    de    fr    es    it
  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 en      0     2     2     7     6     6     6     6
2 no      2     0     1     5     4     6     6     6
3 dk      2     1     0     6     5     6     5     5
4 nl      7     5     6     0     5     9     9     9
5 de      6     4     5     5     0     7     7     7
6 fr      6     6     6     9     7     0     2     1
7 es      6     6     5     9     7     2     0     1
8 it      6     6     5     9     7     1     1     0
9 pl      7     7     6    10     8     5     3     4
10 hu     9     8     8     8     9    10    10    10
11 fi     9     9     9     9     9     9     9     8
# i 3 more variables: pl <dbl>, hu <dbl>, fi <dbl>
```

## Making a distance object

```
number.d %>%
  select(-la) %>%
  as.dist() -> d
d
```

```
      en no dk nl de fr es it pl hu
no    2
dk    2  1
nl    7  5  6
de    6  4  5  5
fr    6  6  6  9  7
```

```

es 6 6 5 9 7 2
it 6 6 5 9 7 1 1
pl 7 7 6 10 8 5 3 4
hu 9 8 8 8 9 10 10 10 10
fi 9 9 9 9 9 9 9 8 9 8

```

```
class(d)
```


```
[1] "dist"
```

## Cluster analysis and dendrogram

```

d.hc <- hclust(d, method = "single")
plot(d.hc)

```



## Comments

- Tree shows how languages combined into clusters.
- First (bottom), Spanish, French, Italian joined into one cluster, Norwegian and Danish into another.
- Later, English joined to Norse languages, Polish to Romance group.
- Then German, Dutch make a Germanic group.
- Finally, Hungarian and Finnish joined to each other and everything else.



## Clustering process

```
d.hc$labels
```

```
[1] "en" "no" "dk" "nl" "de" "fr" "es" "it" "pl" "hu" "fi"
```

```
d.hc$merge
```


	[,1]	[,2]
[1,]	-2	-3
[2,]	-6	-8
[3,]	-7	2
[4,]	-1	1
[5,]	-9	3
[6,]	-5	4
[7,]	-4	6
[8,]	5	7
[9,]	-10	8
[10,]	-11	9

## Comments

- Lines of `merge` show what was combined
  - First, languages 2 and 3 (`no` and `dk`)
  - Then languages 6 and 8 (`fr` and `it`)
  - Then #7 combined with cluster formed at step 2 (`es` joined to `fr` and `it`).
  - Then `en` joined to `no` and `dk` ...
  - Finally `fi` joined to all others.


## Complete linkage

```
d.hc <- hclust(d, method = "complete")  
plot(d.hc)
```



## Ward

```
d.hc <- hclust(d, method = "ward.D")
plot(d.hc)
```



## Chopping the tree

- Three clusters (from Ward) looks good:

```
cutree(d.hc, 3)
```

```
en no dk nl de fr es it pl hu fi
1  1  1  1  1  2  2  2  2  3  3
```

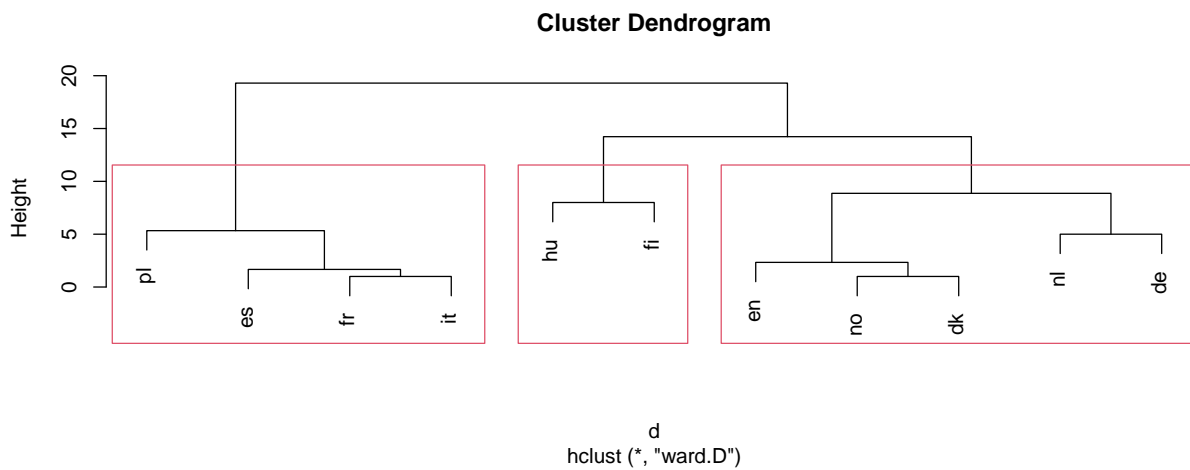
## Turning the “named vector” into a data frame

```
cutree(d.hc, 3) %>% enframe(name="country", value="cluster")
```

```
# A tibble: 11 x 2
  country cluster
  <chr>      <int>
1 en          1
2 no          1
3 dk          1
4 nl          1
5 de          1
6 fr          2
7 es          2
8 it          2
9 pl          2
10 hu         3
11 fi         3
```

## Drawing those clusters on the tree

```
plot(d.hc)
rect.hclust(d.hc, 3)
```



## Comparing single-linkage and Ward

- In Ward, Dutch and German get joined earlier (before joining to Germanic cluster).
- Also Hungarian and Finnish get combined earlier.

## Making those dissimilarities

Original data:

```
my_url <- "http://ritsokiguess.site/datafiles/one-ten.txt"
lang <- read_delim(my_url, " ")
lang
```

```
# A tibble: 10 x 11
  en   no   dk   nl   de   fr   es   it   pl
<chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr>
1 one  en   en   een  eins un   uno  uno  jeden
2 two  to   to   twee zwei  deux dos  due  dwa
3 three tre tre drie drei trois tres tre  trzy
4 four fire fire vier vier quatre cuatro quatt- czte-
5 five fem fem vijf funf cinq cinco cinque piec
6 six seks seks zes sechs six seis sei  szesc
7 seven sju syv zeven sieben sept siete sette sied-
8 eight atte otte acht acht huit ocho otto osiem
9 nine ni   ni   negen neun  neuf nueve nove dzie-
10 ten  ti   ti   tien zehn dix  diez dieci dzie-
# i 2 more variables: hu <chr>, fi <chr>
```

It would be a lot easier to extract the first letter if the number names were all in one column.

## Tidy, and extract first letter

```
lang %>% mutate(number=row_number()) %>%
  pivot_longer(-number, names_to="language", values_to="name") %>%
  mutate(first=str_sub(name, 1, 1)) -> lang.long
lang.long
```

```
# A tibble: 110 x 4
  number language name first
  <int> <chr>    <chr> <chr>
1     1 en      one  o
2     1 no      en   e
3     1 dk      en   e
4     1 nl     een  e
5     1 de     eins e
6     1 fr      un   u
7     1 es     uno  u
8     1 it     uno  u
9     1 pl     jeden j
10    1 hu     egy  e
# i 100 more rows
```

## Calculating dissimilarity

- Suppose we wanted dissimilarity between English and Norwegian. It's the number of first letters that are different.
- First get the lines for English:

```
english <- lang.long %>% filter(language == "en")
english
```

```
# A tibble: 10 x 4
  number language name first
  <int> <chr>    <chr> <chr>
1     1 en      one  o
2     2 en      two  t
3     3 en     three t
4     4 en     four  f
5     5 en     five  f
6     6 en     six   s
7     7 en    seven  s
8     8 en    eight  e
9     9 en     nine  n
10    10 en     ten   t
```

## And then the lines for Norwegian

```
norwegian <- lang.long %>% filter(language == "no")
norwegian
```

```
# A tibble: 10 x 4
  number language name first
  <int> <chr>    <chr> <chr>
1     1 no      en    e
2     2 no      to    t
3     3 no      tre   t
4     4 no      fire  f
5     5 no      fem   f
6     6 no      seks  s
7     7 no      sju   s
8     8 no      atte  a
9     9 no      ni    n
10    10 no      ti    t
```

And now we want to put them side by side, matched by number. This is what `left_join` does. (A “join” is a lookup of values in one table using another.)

## The join

```
english %>% left_join(norwegian, by = "number")
```

```
# A tibble: 10 x 7
  number language.x name.x first.x language.y name.y first.y
  <int> <chr>      <chr> <chr> <chr>      <chr> <chr>
1     1 en        one   o     no        en     e
2     2 en        two   t     no        to     t
3     3 en        three t     no        tre    t
4     4 en        four  f     no        fire   f
5     5 en        five  f     no        fem    f
6     6 en        six   s     no        seks   s
7     7 en        seven s     no        sju    s
8     8 en        eight e     no        atte   a
9     9 en        nine  n     no        ni     n
10    10 en        ten   t     no        ti     t
```

`first.x` is 1st letter of English word, `first.y` 1st letter of Norwegian word.

## Counting the different ones

```
english %>%
  left_join(norwegian, by = "number") %>%
  count(different=(first.x != first.y))
```

```
# A tibble: 2 x 2
  different      n
  <lgl>      <int>
1 FALSE         8
2 TRUE          2
```

or

```
english %>%
  left_join(norwegian, by = "number") %>%
  count(different=(first.x != first.y)) %>%
  filter(different) %>% pull(n) -> ans
ans
```

```
[1] 2
```

Words for 1 and 8 start with different letter; rest are same.

## A language with itself

The answer should be zero:

```
english %>%  
  left_join(english, by = "number") %>%  
  count(different=(first.x != first.y)) %>%  
  filter(different) %>% pull(n) -> ans  
ans
```

integer(0)

- but this is “an integer vector of length zero”.
- so we have to allow for this possibility when we write a function to do it.

## Function to do this for any two languages

```
countdiff <- function(lang.1, lang.2, d) {  
  d %>% filter(language == lang.1) -> lang1d  
  d %>% filter(language == lang.2) -> lang2d  
  lang1d %>%  
    left_join(lang2d, by = "number") %>%  
    count(different = (first.x != first.y)) %>%  
    filter(different) %>% pull(n) -> ans  
  # if ans has length zero, set answer to (integer) zero.  
  ifelse(length(ans)==0, 0L, ans)  
}
```

## Testing

```
countdiff("en", "no", lang.long)
```

```
[1] 2
```

```
countdiff("en", "en", lang.long)
```

```
[1] 0
```

English and Norwegian have two different; English and English have none different.  
Check.

## For all pairs of languages?

- First need all the languages:

```
languages <- names(lang)
languages
```

```
[1] "en" "no" "dk" "nl" "de" "fr" "es" "it" "pl"
[10] "hu" "fi"
```

- and then all *pairs* of languages:

```
pairs <- crossing(lang = languages, lang2 = languages)
```

## The pairs

```
pairs
```

```
# A tibble: 121 x 2
  lang lang2
<chr> <chr>
1 de    de
2 de    dk
3 de    en
4 de    es
5 de    fi
6 de    fr
7 de    hu
8 de    it
9 de    nl
10 de   no
# i 111 more rows
```

## Run countdiff for all those language pairs

```
pairs %>% rowwise() %>%
  mutate(diff = countdiff(lang, lang2, lang.long)) -> thediff
thediff
```



```
# A tibble: 121 x 3
# Rowwise:
  lang lang2 diff
  <chr> <chr> <int>
1 de    de      0
2 de    dk      5
3 de    en      6
4 de    es      7
5 de    fi      9
6 de    fr      7
7 de    hu      9
8 de    it      7
9 de    nl      5
10 de   no      4
# i 111 more rows
```

## Make square table of these

```
thediff %>% pivot_wider(names_from=lang2, values_from=diff)
```

```
# A tibble: 11 x 12
  lang   de   dk   en   es   fi   fr   hu   it
  <chr> <int> <int> <int> <int> <int> <int> <int> <int>
1 de      0    5    6    7    9    7    9    7
2 dk      5    0    2    5    9    6    8    5
3 en      6    2    0    6    9    6    9    6
4 es      7    5    6    0    9    2   10    1
5 fi      9    9    9    9    0    9    8    9
6 fr      7    6    6    2    9    0   10    1
7 hu      9    8    9   10    8   10    0   10
8 it      7    5    6    1    9    1   10    0
9 nl      5    6    7    9    9    9    8    9
10 no      4    1    2    6    9    6    8    6
11 pl      8    6    7    3    9    5   10    4
# i 3 more variables: nl <int>, no <int>, pl <int>
```

and that was where we began.

## Another example

Birth, death and infant mortality rates for 97 countries (variables not dissimilarities):

24.7	5.7	30.8	Albania	12.5	11.9	14.4	Bulgaria
13.4	11.7	11.3	Czechoslovakia	12	12.4	7.6	Former_E._Germany
11.6	13.4	14.8	Hungary	14.3	10.2	16	Poland
13.6	10.7	26.9	Romania	14	9	20.2	Yugoslavia
17.7	10	23	USSR	15.2	9.5	13.1	Byelorussia_SSR

13.4	11.6	13	Ukrainian_SSR	20.7	8.4	25.7	Argentina
46.6	18	111	Bolivia	28.6	7.9	63	Brazil
23.4	5.8	17.1	Chile	27.4	6.1	40	Columbia
32.9	7.4	63	Ecuador	28.3	7.3	56	Guyana
...							

- Want to find groups of similar countries (and how many groups, which countries in each group).
- Tree would be unwieldy with 97 countries.
- More automatic way of finding given number of clusters?

## Reading in

```
url <- "http://ritsokiguess.site/datafiles/birthrate.txt"
vital <- read_table(url)
vital
```

```
# A tibble: 97 x 4
  birth death infant country
  <dbl> <dbl> <dbl> <chr>
1  24.7   5.7   30.8 Albania
2  13.4  11.7   11.3 Czechoslovakia
3  11.6  13.4   14.8 Hungary
4  13.6  10.7   26.9 Romania
5  17.7  10     23    USSR
6  13.4  11.6   13    Ukrainian_SSR
7  46.6  18     111   Bolivia
8  23.4   5.8   17.1 Chile
9  32.9   7.4   63    Ecuador
10 34.8   6.6   42    Paraguay
# i 87 more rows
```

## Standardizing

- Infant mortality rate numbers bigger than others, consequence of measurement scale (arbitrary).
- Standardize (numerical) columns of data frame to have mean 0, SD 1, done by `scale`.

```
vital %>%
  mutate(across(where(is.numeric), \(x) scale(x))) -> vital.s
```

## Three clusters

Pretend we know 3 clusters is good. Take off the column of countries, and run `kmeans` on the resulting data frame, asking for 3 clusters:

```
vital.s %>% select(-country) %>%
  kmeans(3) -> vital.km3
names(vital.km3)
```

```
[1] "cluster"      "centers"      "totss"
[4] "withinss"     "tot.withinss" "betweenss"
[7] "size"         "iter"         "ifault"
```

A lot of output, so look at these individually.

## What's in the output?

- Cluster sizes:

```
vital.km3$size
```

```
[1] 40 25 32
```

- Cluster centres:

```
vital.km3$centers
```

	birth	death	infant
1	-1.0376994	-0.3289046	-0.90669032
2	1.1780071	1.3323130	1.32732200
3	0.3768062	-0.6297388	0.09639258

- Cluster 2 has lower than average rates on everything; cluster 3 has much higher than average.

## Cluster sums of squares and membership

```
vital.km3$withinss
```

```
[1] 17.21617 28.32560 21.53020
```

Cluster 1 compact relative to others (countries in cluster 1 more similar).

```
vital.km3$cluster
```

```
[1] 3 1 1 1 1 1 2 1 3 3 1 2 1 1 1 1 1 1 1 2 2 1 3 3 3 2  
[29] 1 3 1 3 3 1 1 3 3 3 2 2 3 3 2 2 3 2 2 2 3 1 1 1 1 1 3  
[57] 3 3 3 3 1 1 1 1 1 1 1 1 1 3 3 3 3 1 2 1 3 3 2 3 1 3  
[85] 2 2 2 2 3 2 2 2 2 2 3 2 2
```

The cluster membership for each of the 97 countries.

## Store countries and clusters to which they belong

```
vital.3 <- tibble(  
  country = vital.s$country,  
  cluster = vital.km3$cluster  
)
```

Next, which countries in which cluster?

Write function to extract them:

```
get_countries <- function(i, d) {  
  d %>% filter(cluster == i) %>% pull(country)  
}
```

## Cluster membership: cluster 2

```
get_countries(2, vital.3)
```

[1]	"Bolivia"	"Mexico"	"Afghanistan"
[4]	"Iran"	"Bangladesh"	"Gabon"
[7]	"Ghana"	"Namibia"	"Sierra_Leone"
[10]	"Swaziland"	"Uganda"	"Zaire"
[13]	"Cambodia"	"Nepal"	"Angola"
[16]	"Congo"	"Ethiopia"	"Gambia"
[19]	"Malawi"	"Mozambique"	"Nigeria"
[22]	"Somalia"	"Sudan"	"Tanzania"
[25]	"Zambia"		

### Cluster 3

```
get_countries(3, vital.3)
```

[1]	"Albania"	"Ecuador"	"Paraguay"
[4]	"Kuwait"	"Oman"	"Turkey"
[7]	"India"	"Mongolia"	"Pakistan"
[10]	"Algeria"	"Botswana"	"Egypt"
[13]	"Libya"	"Morocco"	"South_Africa"
[16]	"Zimbabwe"	"Brazil"	"Columbia"
[19]	"Guyana"	"Peru"	"Venezuela"
[22]	"Bahrain"	"Iraq"	"Jordan"
[25]	"Lebanon"	"Saudi_Arabia"	"Indonesia"
[28]	"Malaysia"	"Philippines"	"Vietnam"
[31]	"Kenya"	"Tunisia"	

### Cluster 1

```
get_countries(1, vital.3)
```

[1]	"Czechoslovakia"	"Hungary"
[3]	"Romania"	"USSR"
[5]	"Ukrainian_SSR"	"Chile"
[7]	"Uruguay"	"Finland"
[9]	"France"	"Greece"
[11]	"Italy"	"Norway"
[13]	"Spain"	"Switzerland"
[15]	"Austria"	"Canada"
[17]	"Israel"	"China"
[19]	"Korea"	"Singapore"
[21]	"Thailand"	"Bulgaria"

[23]	"Former_E._Germany"	"Poland"
[25]	"Yugoslavia"	"Byelorussia_SSR"
[27]	"Argentina"	"Belgium"
[29]	"Denmark"	"Germany"
[31]	"Ireland"	"Netherlands"
[33]	"Portugal"	"Sweden"
[35]	"U.K."	"Japan"
[37]	"U.S.A."	"United_Arab_Emirates"
[39]	"Hong_Kong"	"Sri_Lanka"

## Problem!

- `kmeans` uses randomization. So result of one run might be different from another run.
- Example: just run again on 3 clusters, table of results:

```
vital.s %>%
  select(-country) %>% kmeans(3) -> vital.km3a
table(
  first = vital.km3$cluster,
  second = vital.km3a$cluster
)
```

	second	1	2	3
first	1	40	0	0
	2	0	24	1
	3	4	0	28

- Clusters are similar but *not same*.

## Solution to this

- `nstart` option on `kmeans` runs that many times, takes best. Should be same every time:

```
vital.s %>%
  select(-country) %>%
  kmeans(3, nstart = 20) -> vital.km3b
```

## How many clusters?

- Three was just a guess.
- Idea: try a whole bunch of `#clusters` (say 2–20), obtain measure of goodness of fit for each, make plot.
- Appropriate measure is `tot.withinss`.
- Run `kmeans` for each `#clusters`, get `tot.withinss` each time.

## Function to get `tot.withinss`

...for an input number of clusters, taking only numeric columns of input data frame:

```
ss <- function(i, d) {  
  d %>%  
    select(where(is.numeric)) %>%  
    kmeans(i, nstart = 20) -> km  
  km$tot.withinss  
}
```

Note: writing function to be as general as possible, so that we can re-use it later.

## Constructing within-cluster SS

Make a data frame with desired numbers of clusters, and fill it with the total within-group sums of squares. `ss` expects a single number of clusters, not a vector of several, so run **rowwise**:

```
tibble(clusters = 2:20) %>%  
  rowwise() %>%  
  mutate(wss = ss(clusters, vital.s)) -> ssd  
ssd
```

```
# A tibble: 19 x 2
```


```
# Rowwise:
```

	clusters	wss
	<int>	<dbl>
1	2	117.
2	3	66.9
3	4	51.4
4	5	37.5

5	6	28.7
6	7	24.7
7	8	22.2
8	9	19.8
9	10	17.5
10	11	16.1
11	12	14.7
12	13	13.3
13	14	12.1
14	15	11.2
15	16	10.5
16	17	9.71
17	18	8.98
18	19	8.44
19	20	7.97

## Scree plot

```
ggplot(ssd, aes(x = clusters, y = wss)) + geom_point() +  
  geom_line()
```



## Interpreting scree plot

- Lower wss better.



- But lower for larger #clusters, harder to explain.
- Compromise: low-ish `wss` and low-ish #clusters.
- Look for “elbow” in plot.
- Idea: this is where `wss` decreases fast then slow.
- On our plot, small elbow at 6 clusters. Try this many clusters.

### Six clusters, using `nstart`

```
set.seed(457299)

vital.s %>%
  select(-country) %>%
  kmeans(6, nstart = 20) -> vital.km6
vital.km6$size
```

```
[1] 24 18 15  2  8 30
```

```
vital.km6$centers
```

	birth	death	infant
1	0.4160993	-0.5169988	0.2648754
2	1.2092406	0.7441347	1.0278003
3	-0.4357690	-1.1438599	-0.7281108
4	-0.2199722	2.1116577	-0.4544435
5	1.3043848	2.1896567	1.9470306
6	-1.1737104	-0.1856375	-0.9534370

### Make a data frame of countries and clusters

```
vital.6 <- tibble(
  country = vital.s$country,
  cluster = vital.km6$cluster
)
vital.6 %>% sample_n(10)
```

```
# A tibble: 10 x 2
  country      cluster
  <chr>        <int>
1 Ghana                2
2 Ukrainian_SSR        6
3 Ethiopia              5
4 Somalia              5
5 Oman                 1
6 Botswana             2
7 Paraguay             1
8 Czechoslovakia        6
9 Peru                 1
10 Afghanistan          5
```

## Cluster 1

Below-average death rate, though other rates a little higher than average:

```
get_countries(1, vital.6)
```

```
[1] "Ecuador"      "Paraguay"      "Oman"
[4] "Turkey"       "India"          "Mongolia"
[7] "Pakistan"     "Algeria"        "Egypt"
[10] "Libya"        "Morocco"        "South_Africa"
[13] "Zimbabwe"     "Brazil"         "Guyana"
[16] "Peru"         "Iraq"           "Jordan"
[19] "Lebanon"      "Saudi_Arabia"   "Indonesia"
[22] "Philippines"  "Vietnam"        "Tunisia"
```

## Cluster 2

High on everything:

```
get_countries(2, vital.6)
```

```
[1] "Bolivia"      "Iran"          "Bangladesh"    "Botswana"
[5] "Gabon"        "Ghana"         "Namibia"       "Swaziland"
[9] "Uganda"       "Zaire"         "Cambodia"      "Nepal"
[13] "Congo"        "Kenya"         "Nigeria"       "Sudan"
[17] "Tanzania"     "Zambia"
```

### Cluster 3

Low on everything:

```
get_countries(3, vital.6)
```

```
[1] "Albania"      "Chile"
[3] "Israel"      "Kuwait"
[5] "China"       "Singapore"
[7] "Thailand"     "Argentina"
[9] "Columbia"    "Venezuela"
[11] "Bahrain"     "United_Arab_Emirates"
[13] "Hong_Kong"   "Malaysia"
[15] "Sri_Lanka"
```

### Cluster 4

Very high death rate, just below average on all else:

```
get_countries(4, vital.6)
```

```
[1] "Mexico" "Korea"
```

### Cluster 5

Very high on everything:

```
get_countries(5, vital.6)
```

```
[1] "Afghanistan" "Sierra_Leone" "Angola"
[4] "Ethiopia"    "Gambia"       "Malawi"
[7] "Mozambique"  "Somalia"
```

### Cluster 6

A bit below average on everything:

```
get_countries(6, vital.6)
```

```

[1] "Czechoslovakia"    "Hungary"
[3] "Romania"           "USSR"
[5] "Ukrainian_SSR"     "Uruguay"
[7] "Finland"           "France"
[9] "Greece"            "Italy"
[11] "Norway"            "Spain"
[13] "Switzerland"       "Austria"
[15] "Canada"            "Bulgaria"
[17] "Former_E._Germany" "Poland"
[19] "Yugoslavia"        "Byelorussia_SSR"
[21] "Belgium"           "Denmark"
[23] "Germany"           "Ireland"
[25] "Netherlands"       "Portugal"
[27] "Sweden"            "U.K."
[29] "Japan"             "U.S.A."

```

## Comparing our 3 and 6-cluster solutions

```
table(three = vital.km3$cluster, six = vital.km6$cluster)
```

```

      six
three 1  2  3  4  5  6
   1  0  0  9  1  0 30
   2  0 16  0  1  8  0
   3 24  2  6  0  0  0

```

Compared to 3-cluster solution:

- most of (old) cluster 1 gone to (new) cluster 6
- cluster 2 split into clusters 2 and 5 (two types of “poor” countries)
- cluster 3 split into clusters 1 and 3 (two types of “intermediate” countries, divided by death rate).

## Getting a picture from kmeans

- Use discriminant analysis on clusters found, treating them as “known” groups.

## Discriminant analysis

- So what makes the groups different?
- Uses package MASS (loaded):

```
vital.lda <- lda(vital.km6$cluster ~ birth + death + infant,  
                data = vital.s)  
vital.lda$svd
```

```
[1] 21.687195  8.851811  1.773006
```

```
vital.lda$scaling
```

	LD1	LD2	LD3
birth	2.6879695	1.1224202	-1.9483853
death	0.6652712	-2.7213044	-0.6049358
infant	2.1111801	0.7650912	2.3542296

- LD1 is some of everything (high=poor, low=rich).
- LD2 mainly death rate, high or low.

## A data frame to make plot from

- Get predictions first:

```
vital.pred <- predict(vital.lda)  
d <- data.frame(  
  country = vital.s$country,  
  cluster = vital.km6$cluster,  
  vital.pred$x  
)  
d
```

	country	cluster	LD1
1	Albania	3	-2.74034473
2	Czechoslovakia	6	-5.01874312
3	Hungary	6	-4.97189595
4	Romania	6	-4.40612396
5	USSR	6	-3.87181416
6	Ukrainian_SSR	6	-4.95502329
7	Bolivia	2	7.04719692

8	Chile	3	-3.61284528
9	Ecuador	1	0.60813286
10	Paraguay	1	-0.09333631
11	Uruguay	6	-3.92003877
12	Mexico	4	1.17794263
13	Finland	6	-5.53992595
14	France	6	-5.48731510
15	Greece	6	-6.04517374
16	Italy	6	-6.23984314
17	Norway	6	-5.14396822
18	Spain	6	-6.20238403
19	Switzerland	6	-5.70503604
20	Austria	6	-5.48811665
21	Canada	6	-5.61852237
22	Afghanistan	5	9.15790453
23	Iran	2	5.17009665
24	Israel	3	-4.09921633
25	Kuwait	3	-3.52238895
26	Oman	1	2.12959716
27	Turkey	1	0.61384977
28	Bangladesh	2	6.18349326
29	China	3	-3.23659422
30	India	1	1.81800150
31	Korea	4	-1.46967404
32	Mongolia	1	1.67300272
33	Pakistan	1	2.24428372
34	Singapore	3	-5.25056428
35	Thailand	3	-3.05879372
36	Algeria	1	1.75779182
37	Botswana	2	4.48834877
38	Egypt	1	1.45515716
39	Gabon	2	5.07955834
40	Ghana	2	4.94529597
41	Libya	1	3.96906505
42	Morocco	1	2.33973207
43	Namibia	2	6.78839862
44	Sierra_Leone	5	10.11147508
45	South_Africa	1	1.22038392
46	Swaziland	2	6.62089564
47	Uganda	2	7.44759122
48	Zaire	2	5.01954548
49	Zimbabwe	1	2.90708292
50	Bulgaria	6	-5.02639603
51	Former_E._Germany	6	-5.36617200
52	Poland	6	-4.83913984
53	Yugoslavia	6	-4.87765104
54	Byelorussia_SSR	6	-4.89387951
55	Argentina	3	-3.38164344

56	Brazil	1	-0.17351090
57	Columbia	3	-1.72504134
58	Guyana	1	-0.64024371
59	Peru	1	2.88979728
60	Venezuela	3	-2.71512137
61	Belgium	6	-5.61006438
62	Denmark	6	-5.36296636
63	Germany	6	-5.66618148
64	Ireland	6	-5.22803491
65	Netherlands	6	-5.69497185
66	Portugal	6	-5.54867384
67	Sweden	6	-5.14801099
68	U.K.	6	-5.14080545
69	Japan	6	-6.74109078
70	U.S.A.	6	-4.98026088
71	Bahrain	3	-2.95751794
72	Iraq	1	2.86550554
73	Jordan	1	0.78337204
74	Lebanon	1	-0.13242340
75	Saudi_Arabia	1	2.82947028
76	United_Arab_Emirates	3	-3.60965788
77	Cambodia	2	6.68714487
78	Hong_Kong	3	-6.56814921
79	Indonesia	1	0.59203984
80	Malaysia	3	-1.69768174
81	Nepal	2	5.98051592
82	Philippines	1	-0.11564367
83	Sri_Lanka	3	-3.86669807
84	Vietnam	1	0.73637772
85	Angola	5	8.67464018
86	Congo	2	4.71698904
87	Ethiopia	5	9.02400482
88	Gambia	5	9.16151584
89	Kenya	2	4.37728422
90	Malawi	5	9.25868839
91	Mozambique	5	8.17837256
92	Nigeria	2	6.80523311
93	Somalia	5	9.02055241
94	Sudan	2	6.19772237
95	Tunisia	1	-0.26827112
96	Tanzania	2	7.01894686
97	Zambia	2	5.90158841
	LD2	LD3	
1	2.23114272	0.086392118	
2	-2.54276395	-0.067491502	
3	-3.62910309	0.149274227	
4	-1.68114304	0.832426535	
5	-0.99643221	0.134219177	

6	-2.45593006	0.032543051
7	-1.82235340	-0.559232315
8	1.83697531	-0.440912763
9	2.45078735	0.333955737
10	2.72730980	-0.910115449
11	-0.75565723	0.086830641
12	-7.45661911	-2.185449656
13	-1.71396053	-0.111993227
14	-1.24432254	0.003489983
15	-1.35732249	0.717192252
16	-1.36850774	0.675127833
17	-1.94087447	-0.245927314
18	-0.77030865	0.612616941
19	-1.39900851	0.133327537
20	0.04445534	0.107554431
21	0.05655923	0.137152032
22	-1.57150002	3.855195867
23	1.59572127	0.728082136
24	1.32996325	-0.726569567
25	4.20168393	-0.538116615
26	2.88622987	-2.722021794
27	1.77493550	1.401386127
28	-0.58998184	0.808514931
29	1.37556784	0.521048584
30	1.07819827	1.747922987
31	-5.22549147	-1.651932383
32	1.97934176	-0.052585403
33	2.56907200	2.904858080
34	1.56461210	-0.048777986
35	0.81462641	0.027926012
36	2.32221009	0.405915880
37	1.35059836	-2.251690523
38	1.48375960	-1.484115076
39	-1.84934413	0.223023990
40	0.51518412	-0.679939548
41	2.51546738	-0.550299598
42	1.57697751	0.620167354
43	1.81616353	1.811177322
44	-4.13640327	0.708806858
45	1.07036311	0.584292394
46	1.53114576	0.486213522
47	-0.08614119	-1.461769491
48	-0.14593127	-1.354022712
49	1.53174980	-1.155638673
50	-2.68287368	0.194600531
51	-3.13019129	-0.146640390
52	-1.51169554	0.238889303
53	-0.76403413	0.653220196



54	-1.07548768	0.052116630
55	0.23391708	0.049204615
56	1.80173696	0.887331423
57	2.37367896	0.116914571
58	2.01175992	0.650267842
59	2.70398544	2.617485939
60	3.09958020	-0.531015134
61	-2.07122472	0.103010698
62	-2.80594142	-0.144208119
63	-2.48058100	0.085615217
64	-0.94271289	-0.168083173
65	-0.81402153	0.149795861
66	-1.34891130	0.526746892
67	-2.19511746	-0.439370072
68	-2.45732616	-0.218666957
69	-0.01816615	0.738650841
70	-0.19798546	-0.186143629
71	3.40403941	-0.956027685
72	3.12008079	-0.806112179
73	3.21739659	-1.371400507
74	1.34062918	-0.430471811
75	3.22903162	-0.605791417
76	3.10639879	0.361277059
77	-1.11737674	1.343455301
78	1.21156621	0.795955999
79	1.12304477	1.306331526
80	2.74828222	-1.241072895
81	-0.24581105	1.734265161
82	2.00054960	-0.669610455
83	1.46702228	-0.073210404
84	1.14664253	0.270008884
85	-2.62832040	0.398976877
86	-0.50507133	-1.989873503
87	-2.80509362	0.132536656
88	-3.21458937	0.521137723
89	1.48515306	-1.740964552
90	-5.46422766	-0.742329819
91	-1.74864060	1.141424236
92	-0.35943688	-0.827233974
93	-2.47121452	-0.274058230
94	-0.74977705	-0.138778600
95	2.17721533	0.042802627
96	0.75977710	-0.855439352
97	0.55264060	-2.233552693


## What's in there; making a plot

- d contains country names, cluster memberships and discriminant scores.
- Plot LD1 against LD2, colouring points by cluster and labelling by country:

```
g <- ggplot(d, aes(  
  x = LD1, y = LD2, colour = factor(cluster),  
  label = country  
) + geom_point() +  
  geom_text_repel(size = 2, max.overlaps = Inf) + guides(colour = "none")
```

## The plot

g



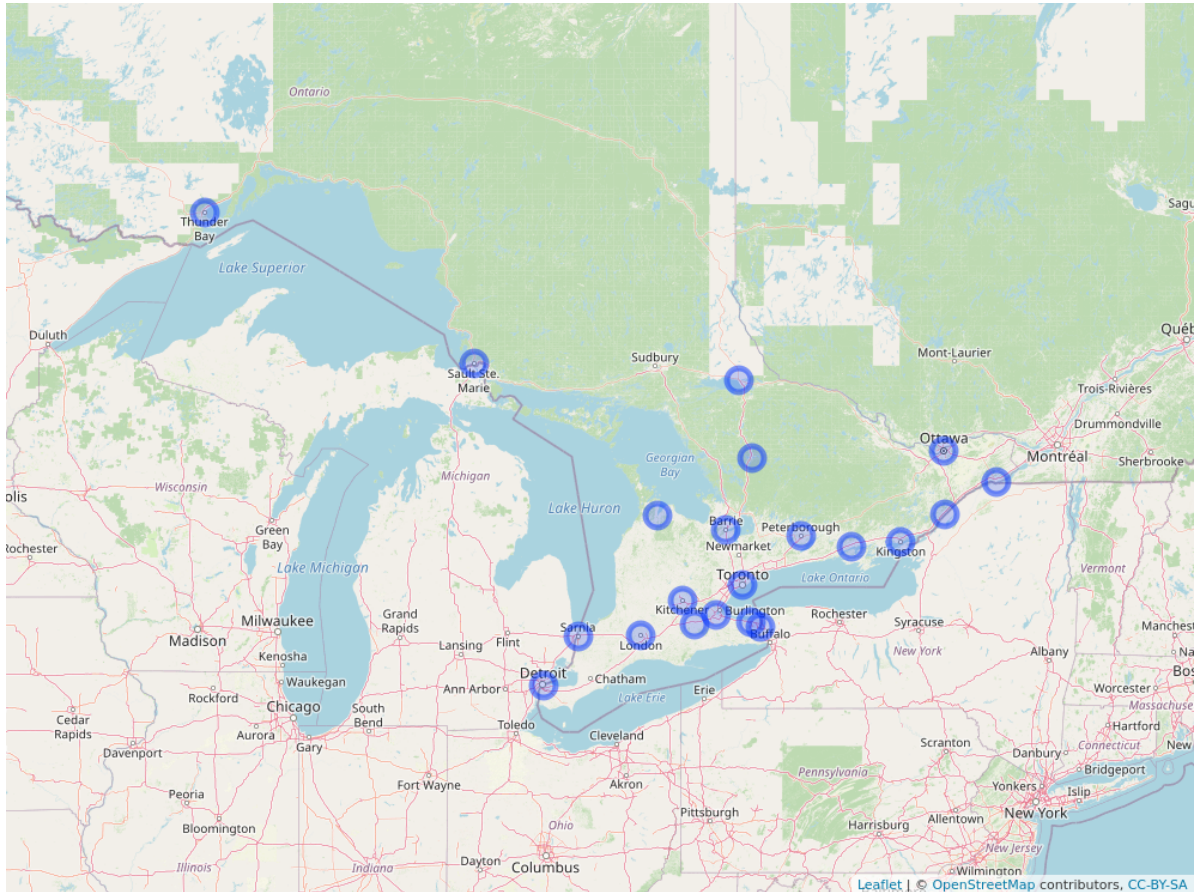
It would be better to zoom in on parts of this plot.

## Final example: a hockey league

- An Ontario hockey league has teams in 21 cities. How can we arrange those teams into 4 geographical divisions?
- Distance data in spreadsheet.

- Take out spaces in team names.
- Save as “text/csv”.
- Distances, so back to `hclust`.

## A map

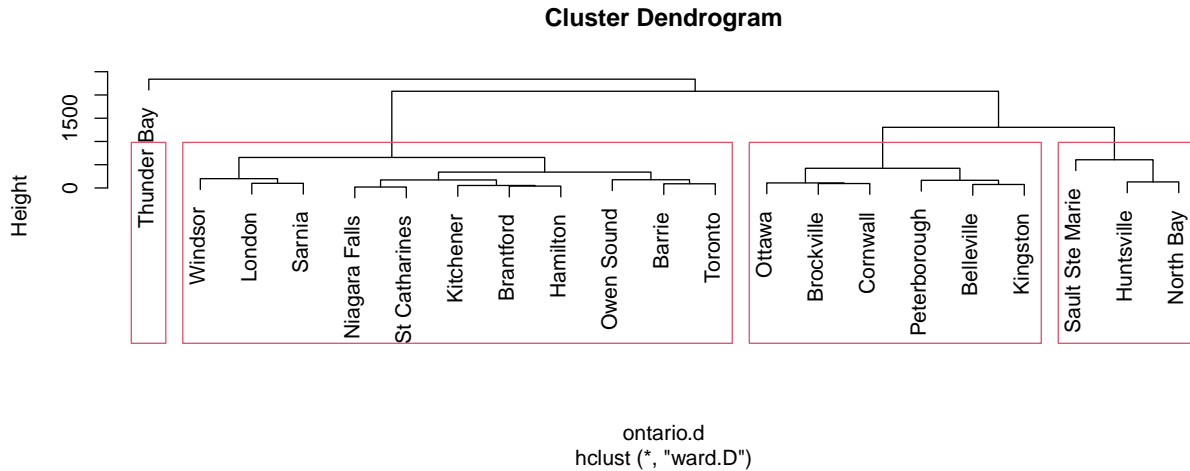


## Attempt 1

```
my_url <-
  "http://ritsokiguess.site/datafiles/ontario-road-distances.csv"
ontario <- read_csv(my_url)
ontario.d <- ontario %>% select(-1) %>% as.dist()
ontario.hc <- hclust(ontario.d, method = "ward.D")
```

## Plot, with 4 clusters

```
plot(ontario.hc)
rect.hclust(ontario.hc, 4)
```

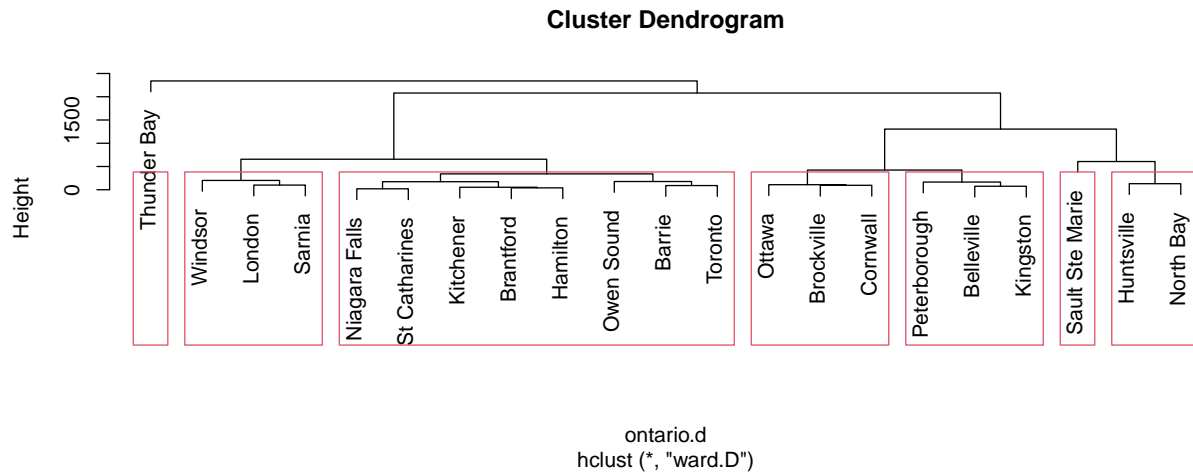


## Comments

- Can't have divisions of 1 team!
- “Southern” divisions way too big!
- Try splitting into more. I found 7 to be good:

## Seven clusters

```
plot(ontario.hc)
rect.hclust(ontario.hc, 7)
```



### Divisions now

- I want to put Huntsville and North Bay together with northern teams.
- I'll put the Eastern teams together. Gives:
- North: Sault Ste Marie, Sudbury, Huntsville, North Bay
- East: Brockville, Cornwall, Ottawa, Peterborough, Belleville, Kingston
- West: Windsor, London, Sarnia
- Central: Owen Sound, Barrie, Toronto, Niagara Falls, St Catharines, Brantford, Hamilton, Kitchener
- Getting them same size beyond us!

## Another map

