

Dates and Times

Packages for this section

```
library(tidyverse)  
# library(lubridate)
```

lubridate is the package that handles dates and times, but is now part of the tidyverse, so no need to load separately.

Dates

- Dates represented on computers as “days since an origin”, typically Jan 1, 1970, with a negative date being before the origin:

```
mydates <- c("1970-01-01", "2007-09-04", "1931-08-05")
(somedates <- tibble(text = mydates) %>%
  mutate(
    d = as.Date(text),
    numbers = as.numeric(d)
  ))
```

```
# A tibble: 3 x 3
```

| | text | d | numbers |
|---|------------|------------|---------|
| | <chr> | <date> | <dbl> |
| 1 | 1970-01-01 | 1970-01-01 | 0 |
| 2 | 2007-09-04 | 2007-09-04 | 13760 |
| 3 | 1931-08-05 | 1931-08-05 | -14029 |

Doing arithmetic with dates

- ▶ Dates are “actually” numbers, so can add and subtract (difference is 2007 date in d minus others):

```
somedates %>% mutate(plus30 = d + 30, diffs = d[2] - d)
```

```
# A tibble: 3 x 5
```

| | text | d | numbers | plus30 | diffs |
|---|------------|------------|---------|------------|------------|
| | <chr> | <date> | <dbl> | <date> | <drtn> |
| 1 | 1970-01-01 | 1970-01-01 | 0 | 1970-01-31 | 13760 days |
| 2 | 2007-09-04 | 2007-09-04 | 13760 | 2007-10-04 | 0 days |
| 3 | 1931-08-05 | 1931-08-05 | -14029 | 1931-09-04 | 27789 days |

Reading in dates from a file

- ▶ `read_csv` and the others can guess that you have dates, if you format them as year-month-day, like column 1 of this `.csv`:

```
date,status,dunno  
2011-08-03,hello,August 3 2011  
2011-11-15,still here,November 15 2011  
2012-02-01,goodbye,February 1 2012
```

- ▶ Then read them in:

```
my_url <- "http://ritsokiguess.site/datafiles/mydates.csv"  
ddd <- read_csv(my_url)
```

- ▶ `read_csv` guessed that the 1st column is dates, but not 3rd.

The data as read in

```
ddd
```

```
# A tibble: 3 x 3
```

| | date | status | dunno |
|--|--------|--------|-------|
| | <date> | <chr> | <chr> |

| | | | |
|---|------------|-------|---------------|
| 1 | 2011-08-03 | hello | August 3 2011 |
|---|------------|-------|---------------|

| | | | |
|---|------------|------------|------------------|
| 2 | 2011-11-15 | still here | November 15 2011 |
|---|------------|------------|------------------|

| | | | |
|---|------------|---------|-----------------|
| 3 | 2012-02-01 | goodbye | February 1 2012 |
|---|------------|---------|-----------------|

Dates in other formats

- ▶ Preceding shows that dates should be stored as text in format yyyy-mm-dd (ISO standard).
- ▶ To deal with dates in other formats, use package lubridate and convert. For example, dates in US format with month first:

```
tibble(usdates = c("05/27/2012", "01/03/2016", "12/31/2015"),  
       mutate(iso = mdy(usdates)))
```

```
# A tibble: 3 x 2
```

| usdates | iso |
|---------|--------|
| <chr> | <date> |

| | |
|--------------|------------|
| 1 05/27/2012 | 2012-05-27 |
| 2 01/03/2016 | 2016-01-03 |
| 3 12/31/2015 | 2015-12-31 |

Trying to read these as UK dates

```
tibble(usdates = c("05/27/2012", "01/03/2016", "12/31/2015"),  
       mutate(uk = dmy(usdates)))
```

```
# A tibble: 3 x 2
```

| | usdates | uk |
|--|---------|--------|
| | <chr> | <date> |

| | | |
|---|------------|----|
| 1 | 05/27/2012 | NA |
|---|------------|----|

| | | |
|---|------------|------------|
| 2 | 01/03/2016 | 2016-03-01 |
|---|------------|------------|

| | | |
|---|------------|----|
| 3 | 12/31/2015 | NA |
|---|------------|----|

- For UK-format dates with month second, one of these dates is legit, but the other two make no sense.

Our data frame's last column:

► Back to this:

```
ddd
```

```
# A tibble: 3 x 3
```

| | date | status | dunno |
|--|--------|--------|-------|
| | <date> | <chr> | <chr> |

| | | | |
|---|------------|-------|---------------|
| 1 | 2011-08-03 | hello | August 3 2011 |
|---|------------|-------|---------------|

| | | | |
|---|------------|------------|------------------|
| 2 | 2011-11-15 | still here | November 15 2011 |
|---|------------|------------|------------------|

| | | | |
|---|------------|---------|-----------------|
| 3 | 2012-02-01 | goodbye | February 1 2012 |
|---|------------|---------|-----------------|

► Month, day, year in that order.

so interpret as such

```
(ddd %>% mutate(date2 = mdy(dunno)) -> d4)
```

```
# A tibble: 3 x 4
```

| | date | status | dunno | date2 |
|---|------------|------------|------------------|------------|
| | <date> | <chr> | <chr> | <date> |
| 1 | 2011-08-03 | hello | August 3 2011 | 2011-08-03 |
| 2 | 2011-11-15 | still here | November 15 2011 | 2011-11-15 |
| 3 | 2012-02-01 | goodbye | February 1 2012 | 2012-02-01 |

Are they really the same?

- ▶ Column date2 was correctly converted from column dunno:

```
d4 %>% mutate(equal = identical(date, date2))
```

```
# A tibble: 3 x 5
```

| | date | status | dunno | date2 | equal |
|---|------------|------------|------------------|------------|-------|
| | <date> | <chr> | <chr> | <date> | <lgl> |
| 1 | 2011-08-03 | hello | August 3 2011 | 2011-08-03 | TRUE |
| 2 | 2011-11-15 | still here | November 15 2011 | 2011-11-15 | TRUE |
| 3 | 2012-02-01 | goodbye | February 1 2012 | 2012-02-01 | TRUE |

- ▶ The two columns of dates are all the same.

Making dates from pieces

Starting from this file:

```
year month day
```

```
1970 1 1
```

```
2007 9 4
```

```
1940 4 15
```

```
my_url <- "http://ritsokiguess.site/datafiles/pieces.txt"  
dates0 <- read_delim(my_url, " ")
```

Making some dates

```
dates0
```

```
# A tibble: 3 x 3
  year month   day
<dbl> <dbl> <dbl>
1  1970     1     1
2  2007     9     4
3  1940     4    15
```

```
dates0 %>%
  unite(dates, day, month, year)%>%
  mutate(d = dmy(dates)) -> newdates
```

The results

```
newdates
```

```
# A tibble: 3 x 2
  dates      d
  <chr>    <date>
1 1_1_1970 1970-01-01
2 4_9_2007 2007-09-04
3 15_4_1940 1940-04-15
```

- ▶ `unite` glues things together with an underscore between them (if you don't specify anything else). Syntax: first thing is new column to be created, other columns are what to make it out of.
- ▶ `unite` makes the original variable columns year, month, day *disappear*.
- ▶ The column `dates` is text, while `d` is a real date.

Extracting information from dates

```
newdates %>%  
  mutate(  
    mon = month(d),  
    day = day(d),  
    weekday = wday(d, label = TRUE)  
  )
```

A tibble: 3 x 5

| | dates | d | mon | day | weekday |
|---|-----------|------------|-------|-------|---------|
| | <chr> | <date> | <dbl> | <int> | <ord> |
| 1 | 1_1_1970 | 1970-01-01 | 1 | 1 | Thu |
| 2 | 4_9_2007 | 2007-09-04 | 9 | 4 | Tue |
| 3 | 15_4_1940 | 1940-04-15 | 4 | 15 | Mon |

Dates and times

- ▶ Standard format for times is to put the time after the date, hours, minutes, seconds:

```
(dd <- tibble(text = c(
  "1970-01-01 07:50:01", "2007-09-04 15:30:00",
  "1940-04-15 06:45:10", "2016-02-10 12:26:40"
)))
```

```
# A tibble: 4 x 1
```

```
text
```

```
<chr>
```

```
1 1970-01-01 07:50:01
2 2007-09-04 15:30:00
3 1940-04-15 06:45:10
4 2016-02-10 12:26:40
```


Converting text to date-times:

- ▶ Then get from this text using `ymd_hms`:

```
dd %>% mutate(dt = ymd_hms(text))
```

```
# A tibble: 4 x 2
```

| | text | dt |
|---|---------------------|---------------------|
| | <chr> | <dtm> |
| 1 | 1970-01-01 07:50:01 | 1970-01-01 07:50:01 |
| 2 | 2007-09-04 15:30:00 | 2007-09-04 15:30:00 |
| 3 | 1940-04-15 06:45:10 | 1940-04-15 06:45:10 |
| 4 | 2016-02-10 12:26:40 | 2016-02-10 12:26:40 |

Timezones

- ▶ Default timezone is “Universal Coordinated Time”. Change it via `tz=` and the name of a timezone:

```
dd %>%  
  mutate(dt = ymd_hms(text, tz = "America/Toronto")) -> dd  
dd %>% mutate(zone = tz(dt))
```

A tibble: 4 x 3

| | text | dt | zone |
|---|---------------------|---------------------|-----------------|
| | <chr> | <dtm> | <chr> |
| 1 | 1970-01-01 07:50:01 | 1970-01-01 07:50:01 | America/Toronto |
| 2 | 2007-09-04 15:30:00 | 2007-09-04 15:30:00 | America/Toronto |
| 3 | 1940-04-15 06:45:10 | 1940-04-15 06:45:10 | America/Toronto |
| 4 | 2016-02-10 12:26:40 | 2016-02-10 12:26:40 | America/Toronto |

Extracting time parts

► As you would expect:

```
dd %>%  
  select(-text) %>%  
  mutate(  
    h = hour(dt),  
    sec = second(dt),  
    min = minute(dt),  
    zone = tz(dt)  
  )
```

A tibble: 4 x 5

| | dt | h | sec | min | zone |
|---|---------------------|-------|-------|-------|-----------------|
| | <dtm> | <int> | <dbl> | <int> | <chr> |
| 1 | 1970-01-01 07:50:01 | 7 | 1 | 50 | America/Toronto |
| 2 | 2007-09-04 15:30:00 | 15 | 0 | 30 | America/Toronto |
| 3 | 1940-04-15 06:45:10 | 6 | 10 | 45 | America/Toronto |
| 4 | 2016-02-10 12:26:40 | 12 | 40 | 26 | America/Toronto |

Same times, but different time zone:

```
dd %>%  
  select(dt) %>%  
  mutate(oz = with_tz(dt, "Australia/Sydney"))
```

A tibble: 4 x 2

| | dt | oz |
|---|---------------------|---------------------|
| | <dtm> | <dtm> |
| 1 | 1970-01-01 07:50:01 | 1970-01-01 22:50:01 |
| 2 | 2007-09-04 15:30:00 | 2007-09-05 05:30:00 |
| 3 | 1940-04-15 06:45:10 | 1940-04-15 21:45:10 |
| 4 | 2016-02-10 12:26:40 | 2016-02-11 04:26:40 |

In more detail:

```
dd %>%  
  mutate(oz = with_tz(dt, "Australia/Sydney")) %>%  
  pull(oz)
```

```
[1] "1970-01-01 22:50:01 AEST" "2007-09-05 05:30:00 AEST"  
[3] "1940-04-15 21:45:10 AEST" "2016-02-11 04:26:40 AEDT"
```

How long between date-times?

- ▶ We may need to calculate the time between two events. For example, these are the dates and times that some patients were admitted to and discharged from a hospital:

`admit,discharge`

`1981-12-10 22:00:00,1982-01-03 14:00:00`

`2014-03-07 14:00:00,2014-03-08 09:30:00`

`2016-08-31 21:00:00,2016-09-02 17:00:00`

Do they get read in as date-times?

- These ought to get read in and converted to date-times:

```
my_url <- "http://ritsokiguess.site/datafiles/hospital.csv"
stays <- read_csv(my_url)
stays
```

```
# A tibble: 3 x 2
```

| | admit <dtm> | discharge <dtm> |
|---|---------------------|---------------------|
| 1 | 1981-12-10 22:00:00 | 1982-01-03 14:00:00 |
| 2 | 2014-03-07 14:00:00 | 2014-03-08 09:30:00 |
| 3 | 2016-08-31 21:00:00 | 2016-09-02 17:00:00 |

- and so it proves.

Subtracting the date-times

► In the obvious way, this gets us an answer:

```
stays %>% mutate(stay = discharge - admit)
```

```
# A tibble: 3 x 3
```

| | admit <dtm> | discharge <dtm> | stay <drtn> |
|---|---------------------|---------------------|----------------|
| 1 | 1981-12-10 22:00:00 | 1982-01-03 14:00:00 | 568.0 hours |
| 2 | 2014-03-07 14:00:00 | 2014-03-08 09:30:00 | 19.5 hours |
| 3 | 2016-08-31 21:00:00 | 2016-09-02 17:00:00 | 44.0 hours |

► Number of hours; hard to interpret.

Days

- Fractional number of days would be better:

```
stays %>%  
  mutate(  
    stay_days = as.period(admit %--% discharge) / days(1))
```

A tibble: 3 x 3

| | admit <dtm> | discharge <dtm> | stay_days <dbl> |
|---|---------------------|---------------------|--------------------|
| 1 | 1981-12-10 22:00:00 | 1982-01-03 14:00:00 | 23.7 |
| 2 | 2014-03-07 14:00:00 | 2014-03-08 09:30:00 | 0.812 |
| 3 | 2016-08-31 21:00:00 | 2016-09-02 17:00:00 | 1.83 |

Completed days

- Pull out with `day()` etc, as for a date-time:

```
stays %>%  
  mutate(  
    stay = as.period(admit %--% discharge),  
    stay_days = day(stay),  
    stay_hours = hour(stay)  
  ) %>%  
  select(starts_with("stay"))
```

A tibble: 3 x 3

| | stay | stay_days | stay_hours |
|---|---------------|-----------|------------|
| | <Period> | <dbl> | <dbl> |
| 1 | 23d 16H 0M 0S | 23 | 16 |
| 2 | 19H 30M 0S | 0 | 19 |
| 3 | 1d 20H 0M 0S | 1 | 20 |

Comments

- ▶ Date-times are stored internally as seconds-since-something, so that subtracting two of them will give, internally, a number of seconds.
- ▶ Just subtracting the date-times is displayed as a time (in units that R chooses for us).
- ▶ Convert to fractional times via a “period”, then divide by `days(1)`, `months(1)` etc.
- ▶ These ideas useful for calculating time from a start point until an event happens (in this case, a patient being discharged from hospital).