

Dates and Times

Packages for this section

```
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.2      v readr      2.1.4
v forcats    0.5.0      v stringr    1.5.0
v ggplot2    3.4.2      v tibble     3.2.1
v lubridate  1.9.2      v tidyr      1.3.0
v purrr      1.0.1
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

```
# library(lubridate)
```

`lubridate` is the package that handles dates and times, but is now part of the `tidyverse`, so no need to load separately.

Dates

- Dates represented on computers as “days since an origin”, typically Jan 1, 1970, with a negative date being before the origin:

```
mydates <- c("1970-01-01", "2007-09-04", "1931-08-05")
(somedates <- tibble(text = mydates) %>%
  mutate(
    d = as.Date(text),
```

```

    numbers = as.numeric(d)
  ))

```

```

# A tibble: 3 x 3
  text      d      numbers
  <chr>    <date>    <dbl>
1 1970-01-01 1970-01-01      0
2 2007-09-04 2007-09-04  13760
3 1931-08-05 1931-08-05 -14029

```

Doing arithmetic with dates

- Dates are “actually” numbers, so can add and subtract (difference is 2007 date in `d` minus others):

```

somedates %>% mutate(plus30 = d + 30, diffs = d[2] - d)

```

```

# A tibble: 3 x 5
  text      d      numbers plus30      diffs
  <chr>    <date>    <dbl> <date>    <drtn>
1 1970-01-01 1970-01-01      0 1970-01-31 13760 days
2 2007-09-04 2007-09-04  13760 2007-10-04      0 days
3 1931-08-05 1931-08-05 -14029 1931-09-04 27789 days

```

Reading in dates from a file

- `read_csv` and the others can guess that you have dates, if you format them as year-month-day, like column 1 of this `.csv`:

```

date,status,dunno
2011-08-03,hello,August 3 2011
2011-11-15,still here,November 15 2011
2012-02-01,goodbye,February 1 2012

```

- Then read them in:

```

my_url <- "http://ritsokiguess.site/datafiles/mydates.csv"
ddd <- read_csv(my_url)

```

```

Rows: 3 Columns: 3
-- Column specification -----
Delimiter: ","
chr (2): status, dunno
date (1): date

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

• read_csv guessed that the 1st column is dates, but not 3rd.

```

The data as read in

```

ddd

# A tibble: 3 x 3
  date      status  dunno
  <date>    <chr>    <chr>
1 2011-08-03 hello    August 3 2011
2 2011-11-15 still here November 15 2011
3 2012-02-01 goodbye   February 1 2012

```

Dates in other formats

- Preceding shows that dates should be stored as text in format yyyy-mm-dd (ISO standard).
- To deal with dates in other formats, use package `lubridate` and `convert`. For example, dates in US format with month first:

```

tibble(usdates = c("05/27/2012", "01/03/2016", "12/31/2015")) %>%
  mutate(iso = mdy(usdates))

# A tibble: 3 x 2
  usdates    iso
  <chr>      <date>
1 05/27/2012 2012-05-27
2 01/03/2016 2016-01-03
3 12/31/2015 2015-12-31

```

Trying to read these as UK dates

```
tibble(usdates = c("05/27/2012", "01/03/2016", "12/31/2015")) %>%  
  mutate(uk = dmy(usdates))
```

```
Warning: There was 1 warning in `mutate()`.  
i In argument: `uk = dmy(usdates)`.  
Caused by warning:  
! 2 failed to parse.
```

```
# A tibble: 3 x 2  
  usdates      uk  
  <chr>      <date>  
1 05/27/2012 NA  
2 01/03/2016 2016-03-01  
3 12/31/2015 NA
```

- For UK-format dates with month second, one of these dates is legit, but the other two make no sense.

Our data frame's last column:

- Back to this:

```
ddd
```

```
# A tibble: 3 x 3  
  date      status    dunno  
  <date>    <chr>      <chr>  
1 2011-08-03 hello      August 3 2011  
2 2011-11-15 still here November 15 2011  
3 2012-02-01 goodbye    February 1 2012
```

- Month, day, year in that order.

so interpret as such

```
(ddd %>% mutate(date2 = mdy(dunno)) -> d4)
```

```
# A tibble: 3 x 4
  date      status      dunno      date2
<date>    <chr>    <chr>    <date>
1 2011-08-03 hello      August 3 2011 2011-08-03
2 2011-11-15 still here November 15 2011 2011-11-15
3 2012-02-01 goodbye     February 1 2012 2012-02-01
```

Are they really the same?

- Column date2 was correctly converted from column dunno:

```
d4 %>% mutate(equal = identical(date, date2))
```

```
# A tibble: 3 x 5
  date      status      dunno      date2      equal
<date>    <chr>    <chr>    <date>    <lgl>
1 2011-08-03 hello      August 3 2011 2011-08-03 TRUE
2 2011-11-15 still here November 15 2011 2011-11-15 TRUE
3 2012-02-01 goodbye     February 1 2012 2012-02-01 TRUE
```

- The two columns of dates are all the same.

Making dates from pieces

Starting from this file:

```
year month day
1970 1 1
2007 9 4
1940 4 15
```

```
my_url <- "http://ritsokiguess.site/datafiles/pieces.txt"
dates0 <- read_delim(my_url, " ")
```

```

Rows: 3 Columns: 3
-- Column specification -----
Delimiter: " "
dbl (3): year, month, day

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```

Making some dates

```
dates0
```

```

# A tibble: 3 x 3
  year month  day
<dbl> <dbl> <dbl>
1  1970     1     1
2  2007     9     4
3  1940     4    15

```

```

dates0 %>%
  unite(dates, day, month, year) %>%
  mutate(d = dmy(dates)) -> newdates

```

The results

```
newdates
```

```

# A tibble: 3 x 2
  dates      d
<chr>    <date>
1 1_1_1970 1970-01-01
2 4_9_2007 2007-09-04
3 15_4_1940 1940-04-15

```

- `unite` glues things together with an underscore between them (if you don't specify anything else). Syntax: first thing is new column to be created, other columns are what to make it out of.
- `unite` makes the original variable columns `year`, `month`, `day` *disappear*.
- The column `dates` is text, while `d` is a real date.

Extracting information from dates

```
newdates %>%  
  mutate(  
    mon = month(d),  
    day = day(d),  
    weekday = wday(d, label = T)  
  )
```

```
# A tibble: 3 x 5  
  dates      d      mon    day weekday  
  <chr>    <date>    <dbl> <int> <ord>  
1 1_1_1970 1970-01-01     1     1 Thu  
2 4_9_2007 2007-09-04     9     4 Tue  
3 15_4_1940 1940-04-15     4    15 Mon
```

Dates and times

- Standard format for times is to put the time after the date, hours, minutes, seconds:

```
(dd <- tibble(text = c(  
  "1970-01-01 07:50:01", "2007-09-04 15:30:00",  
  "1940-04-15 06:45:10", "2016-02-10 12:26:40"  
)))
```

```
# A tibble: 4 x 1  
  text  
  <chr>  
1 1970-01-01 07:50:01  
2 2007-09-04 15:30:00  
3 1940-04-15 06:45:10  
4 2016-02-10 12:26:40
```

Converting text to date-times:

- Then get from this text using `ymd_hms`:

```
dd %>% mutate(dt = ymd_hms(text))
```

```
# A tibble: 4 x 2
  text          dt
  <chr>        <dtm>
1 1970-01-01 07:50:01 1970-01-01 07:50:01
2 2007-09-04 15:30:00 2007-09-04 15:30:00
3 1940-04-15 06:45:10 1940-04-15 06:45:10
4 2016-02-10 12:26:40 2016-02-10 12:26:40
```

Timezones

- Default timezone is “Universal Coordinated Time”. Change it via `tz=` and the name of a timezone:

```
dd %>%
  mutate(dt = ymd_hms(text, tz = "America/Toronto")) -> dd
dd %>% mutate(zone = tz(dt))
```

```
# A tibble: 4 x 3
  text          dt          zone
  <chr>        <dtm>        <chr>
1 1970-01-01 07:50:01 1970-01-01 07:50:01 America/Toronto
2 2007-09-04 15:30:00 2007-09-04 15:30:00 America/Toronto
3 1940-04-15 06:45:10 1940-04-15 06:45:10 America/Toronto
4 2016-02-10 12:26:40 2016-02-10 12:26:40 America/Toronto
```

Extracting time parts

- As you would expect:

```
dd %>%
  select(-text) %>%
  mutate(
    h = hour(dt),
    sec = second(dt),
    min = minute(dt),
    zone = tz(dt)
  )
```

```
# A tibble: 4 x 5
  dt          h    sec    min zone
  <dtm>      <dbl> <dbl> <dbl> <chr>
```


	<dtm>	<int>	<dbl>	<int>	<chr>
1	1970-01-01 07:50:01	7	1	50	America/Toronto
2	2007-09-04 15:30:00	15	0	30	America/Toronto
3	1940-04-15 06:45:10	6	10	45	America/Toronto
4	2016-02-10 12:26:40	12	40	26	America/Toronto

Same times, but different time zone:

```
dd %>%
  select(dt) %>%
  mutate(oz = with_tz(dt, "Australia/Sydney"))
```

```
# A tibble: 4 x 2
  dt          oz
<dtm>      <dtm>
1 1970-01-01 07:50:01 1970-01-01 22:50:01
2 2007-09-04 15:30:00 2007-09-05 05:30:00
3 1940-04-15 06:45:10 1940-04-15 21:45:10
4 2016-02-10 12:26:40 2016-02-11 04:26:40
```

In more detail:

```
dd %>%
  mutate(oz = with_tz(dt, "Australia/Sydney")) %>%
  pull(oz)
```

```
[1] "1970-01-01 22:50:01 AEST" "2007-09-05 05:30:00 AEST"
[3] "1940-04-15 21:45:10 AEST" "2016-02-11 04:26:40 AEDT"
```

How long between date-times?

- We may need to calculate the time between two events. For example, these are the dates and times that some patients were admitted to and discharged from a hospital:

```
admit,discharge
1981-12-10 22:00:00,1982-01-03 14:00:00
2014-03-07 14:00:00,2014-03-08 09:30:00
2016-08-31 21:00:00,2016-09-02 17:00:00
```

Do they get read in as date-times?

- These ought to get read in and converted to date-times:

```
my_url <- "http://ritsokiguess.site/datafiles/hospital.csv"
stays <- read_csv(my_url)
```

Rows: 3 Columns: 2

-- Column specification -----

Delimiter: ","

dtm (2): admit, discharge

- i Use ``spec()`` to retrieve the full column specification for this data.
- i Specify the column types or set ``show_col_types = FALSE`` to quiet this message.

```
stays
```

A tibble: 3 x 2

	admit <dtm>	discharge <dtm>
1	1981-12-10 22:00:00	1982-01-03 14:00:00
2	2014-03-07 14:00:00	2014-03-08 09:30:00
3	2016-08-31 21:00:00	2016-09-02 17:00:00

- and so it proves.

Subtracting the date-times

- In the obvious way, this gets us an answer:

```
stays %>% mutate(stay = discharge - admit)
```

A tibble: 3 x 3

	admit <dtm>	discharge <dtm>	stay <drtn>
1	1981-12-10 22:00:00	1982-01-03 14:00:00	568.0 hours
2	2014-03-07 14:00:00	2014-03-08 09:30:00	19.5 hours
3	2016-08-31 21:00:00	2016-09-02 17:00:00	44.0 hours

- Number of hours; hard to interpret.

Days

- Fractional number of days would be better:

```
stays %>%  
  mutate(  
    stay_days = as.period(admit %--% discharge) / days(1))
```

```
# A tibble: 3 x 3  
  admit          discharge      stay_days  
  <dtm>          <dtm>          <dbl>  
1 1981-12-10 22:00:00 1982-01-03 14:00:00    23.7  
2 2014-03-07 14:00:00 2014-03-08 09:30:00     0.812  
3 2016-08-31 21:00:00 2016-09-02 17:00:00     1.83
```

Completed days

- Pull out with `day()` etc, as for a date-time:

```
stays %>%  
  mutate(  
    stay = as.period(admit %--% discharge),  
    stay_days = day(stay),  
    stay_hours = hour(stay)  
  ) %>%  
  select(starts_with("stay"))
```

```
# A tibble: 3 x 3  
  stay          stay_days stay_hours  
  <Period>      <dbl>      <dbl>  
1 23d 16H 0M 0S      23          16  
2 19H 30M 0S         0          19  
3 1d 20H 0M 0S       1          20
```

Comments

- Date-times are stored internally as seconds-since-something, so that subtracting two of them will give, internally, a number of seconds.
- Just subtracting the date-times is displayed as a time (in units that R chooses for us).

- Convert to fractional times via a “period”, then divide by `days(1)`, `months(1)` etc.
- These ideas useful for calculating time from a start point until an event happens (in this case, a patient being discharged from hospital).