

Multivariate Analysis of Variance

Multivariate analysis of variance

- Standard ANOVA has just one response variable.
- What if you have more than one response?
- Try an ANOVA on each response separately.
- But might miss some kinds of interesting dependence between the responses that distinguish the groups.

Packages

```
library(car) # may need to install first
library(tidyverse)
library(MVTests) # also may need to install
```

Small example

- Measure yield and seed weight of plants grown under 2 conditions: low and high amounts of fertilizer.
- Data (fertilizer, yield, seed weight):

```
url <- "http://ritsokiguess.site/datafiles/manova1.txt"
hilo <- read_delim(url, " ")
```

- 2 responses, yield and seed weight.

The data

```
hilo
```

```
# A tibble: 8 x 3
  fertilizer yield weight
  <chr>      <dbl> <dbl>
1 low        34     10
2 low        29     14
3 low        35     11
4 low        32     13
5 high       33     14
6 high       38     12
7 high       34     13
8 high       35     14
```

Boxplot for yield for each fertilizer group


```
ggplot(hilo, aes(x = fertilizer, y = yield)) + geom_boxplot()
```



Yields overlap for fertilizer groups.

Boxplot for weight for each fertilizer group

```
ggplot(hilo, aes(x = fertilizer, y = weight)) + geom_boxplot()
```



Weights overlap for fertilizer groups.

ANOVAs for yield and weight

```
hilo.y <- aov(yield ~ fertilizer, data = hilo)
summary(hilo.y)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
fertilizer	1	12.5	12.500	2.143	0.194
Residuals	6	35.0	5.833		

```
hilo.w <- aov(weight ~ fertilizer, data = hilo)
summary(hilo.w)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
fertilizer	1	3.125	3.125	1.471	0.271
Residuals	6	12.750	2.125		

Neither response depends significantly on fertilizer. But...

Plotting both responses at once

- Have two response variables (not more), so can plot the response variables against *each other*, labelling points by which fertilizer group they're from.
- First, create data frame with points (31, 14) and (38, 10) (why? Later):


```
d <- tribble(
  ~line_x, ~line_y,
  31, 14,
```

```
38, 10
)
```

- Then plot data as points, and add line through points in d:

```
ggplot(hilo, aes(x = yield, y = weight,
                 colour = fertilizer)) + geom_point() +
  geom_line(data = d,
            aes(x = line_x, y = line_y, colour = NULL)) -> g
```

The plot



Comments

- Graph construction:
 - Joining points in d by line.
 - `geom_line` inherits `colour` from `aes` in `ggplot`.
 - Data frame d has no `fertilizer` (previous `colour`), so have to unset.
- Results:
 - High-fertilizer plants have both yield and weight high.
 - True even though no sig difference in yield or weight individually.
 - Drew line separating highs from lows on plot.

MANOVA finds multivariate differences

- Is difference found by diagonal line significant? MANOVA finds out.

```
response <- with(hilo, cbind(yield, weight))
hilo.1 <- manova(response ~ fertilizer, data = hilo)
summary(hilo.1)
```

```
          Df  Pillai approx F num Df den Df  Pr(>F)
fertilizer 1 0.80154   10.097      2     5 0.01755 *
Residuals  6
---
Signif. codes:
0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- Yes! Difference between groups is *diagonally*, not just up/down (weight) or left-right (yield). The *yield-weight combination* matters.

Strategy

- Create new response variable by gluing together columns of responses, using `cbind`.
- Use `manova` with new response, looks like `lm` otherwise.
- With more than 2 responses, cannot draw graph. What then?
- If MANOVA test significant, cannot use Tukey. What then?
- Use *discriminant analysis* (of which more later).

Another way to do MANOVA

using `Manova` from package `car`:

```
hilo.2.lm <- lm(response ~ fertilizer, data = hilo)
hilo.2 <- Manova(hilo.2.lm)
summary(hilo.2)
```

Type II MANOVA Tests:

Sum of squares and products for error:

```
      yield weight
yield   35 -18.00
weight -18  12.75
```

Term: fertilizer

Sum of squares and products for the hypothesis:

```
      yield weight
```

```

yield 12.50 6.250
weight 6.25 3.125

Multivariate Tests: fertilizer
          Df test stat approx F num Df den Df Pr(>F)
Pillai    1  0.801542 10.09714      2    5 0.017546 *
Wilks     1  0.198458 10.09714      2    5 0.017546 *
Hotelling-Lawley 1  4.038855 10.09714      2    5 0.017546 *
Roy       1  4.038855 10.09714      2    5 0.017546 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Comments

- Same result as small-m `manova`.
- Manova will also do *repeated measures*, coming up later.

Assumptions

- normality of each response variable within each treatment group
 - this is actually *multivariate* normality, with correlations
- equal spreads: each response variable has same variances and correlations (with other response variables) within each treatment group. Here:
 - yield has same spread for low and high fertilizer
 - weight has same spread for low and high fertilizer
 - correlation between yield and weight is same for low and high fertilizer
- test equal spread using Box’s *M* test
 - a certain amount of unequalness is OK, so only a concern if P-value from *M*-test is very small (eg. less than 0.001).

Assumptions for yield-weight data

For normal quantile plots, need “extra-long” with all the data values in one column:


```

hilo %>%
  pivot_longer(-fertilizer, names_to = "xname",
               values_to = "xvalue") %>%
  ggplot(aes(sample = xvalue)) + stat_qq() +
  stat_qq_line() +
  facet_grid(xname ~ fertilizer, scales = "free") -> g

```

There are only four observations per response variable - treatment group combination, so graphs are not very informative (over):

The plots



Box M test

- Make sure package `MVTests` loaded first.
- inputs:
 - the response matrix (or, equivalently, the response-variable columns from your dataframe)
 - the column with the grouping variable in it (most easily gotten with `$`).

```
hilo %>% select(yield, weight) -> numeric_values  
summary(BoxM(numeric_values, hilo$fertilizer))
```

Box's M Test

Chi-Squared Value = 1.002964 , df = 3 and p-value: 0.801

- No problem at all with unequal spreads.

Another example: peanuts

- Three different varieties of peanuts (mysteriously, 5, 6 and 8) planted in two different locations.

- Three response variables: `y`, `smk` and `w`.

```
u <- "http://ritsokiguess.site/datafiles/peanuts.txt"
peanuts.orig <- read_delim(u, " ")
```

The data

```
peanuts.orig
```

```
# A tibble: 12 x 6
  obs location variety     y     smk     w
  <dbl>   <dbl>   <dbl> <dbl> <dbl> <dbl>
1     1     1     1     5  195.  153.  51.4
2     2     1     1     5  194.  168.  53.7
3     3     2     2     5  190.  140.  55.5
4     4     2     2     5  180.  121.  44.4
5     5     1     6    203  157.  49.8
6     6     1     6    196  166   45.8
7     7     2     6    203  166   60.4
8     8     2     6    198  162   54.1
9     9     1     8    194  164   57.8
10    10     1     8    187  165   58.6
11    11     2     8    202  167    65
12    12     2     8    200  174   67.2
```

Setup for analysis

```
peanuts.orig %>%
  mutate(
    location = factor(location),
    variety = factor(variety)
  ) -> peanuts
response <- with(peanuts, cbind(y, smk, w))
head(response)
```

```
      y     smk     w
[1,] 195.3 153.1 51.4
[2,] 194.3 167.7 53.7
[3,] 189.7 139.5 55.5
[4,] 180.4 121.1 44.4
```



```
[5,] 203.0 156.8 49.8
[6,] 195.9 166.0 45.8
```

Analysis (using Manova)

```
peanuts.1 <- lm(response ~ location * variety, data = peanuts)
peanuts.2 <- Manova(peanuts.1)
summary(peanuts.2)
```

Type II MANOVA Tests:

Sum of squares and products for error:

	y	smk	w
y	104.205	49.365	76.480
smk	49.365	352.105	121.995
w	76.480	121.995	94.835

Term: location

Sum of squares and products for the hypothesis:

	y	smk	w
y	0.7008333	-10.6575	7.129167
smk	-10.6575000	162.0675	-108.412500
w	7.1291667	-108.4125	72.520833

Multivariate Tests: location

	Df	test stat	approx F	num Df	den Df	Pr(>F)
Pillai	1	0.893484	11.18432	3	4	0.020502 *
Wilks	1	0.106516	11.18432	3	4	0.020502 *
Hotelling-Lawley	1	8.388243	11.18432	3	4	0.020502 *
Roy	1	8.388243	11.18432	3	4	0.020502 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Term: variety

Sum of squares and products for the hypothesis:

	y	smk	w
y	196.1150	365.1825	42.6275
smk	365.1825	1089.0150	414.6550
w	42.6275	414.6550	284.1017

Multivariate Tests: variety

	Df	test stat	approx F	num Df	den Df	Pr(>F)
Pillai	2	1.709109	9.792388	6	10	0.0010562 **
Wilks	2	0.012444	10.619086	6	8	0.0019275 **
Hotelling-Lawley	2	21.375675	10.687838	6	6	0.0054869 **
Roy	2	18.187611	30.312685	3	5	0.0012395 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Term: location:variety

Sum of squares and products for the hypothesis:

	y	smk	w
y	205.1017	363.6675	107.78583
smk	363.6675	780.6950	254.22000
w	107.7858	254.2200	85.95167

Multivariate Tests: location:variety

	Df	test stat	approx F	num Df	den Df	Pr(>F)
Pillai	2	1.290861	3.033867	6	10	0.058708 .
Wilks	2	0.074300	3.558197	6	8	0.050794 .
Hotelling-Lawley	2	7.544290	3.772145	6	6	0.065517 .
Roy	2	6.824094	11.373490	3	5	0.011340 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Comments

- Interaction not quite significant, but main effects are.
- Combined response variable (y, smk, w) definitely depends on location and on variety
- Weak dependence of (y, smk, w) on the location-variety *combination*.
- Understanding that dependence beyond our scope right now.

Comments

- this time there are only six observations per location and four per variety, so normality is still difficult to be confident about
- y at location 1 seems to be the worst for normality (long tails / outliers), and maybe y at location 2 is skewed left, but the others are not bad
- there is some evidence of unequal spread (slopes of lines), but is it bad enough to worry about? (Box M-test, over).

Box's M tests

- One for location, one for variety:

```
summary(BoxM(response, peanuts$location))
```

Box's M Test

Chi-Squared Value = 12.47797 , df = 6 and p-value: 0.0521

```
summary(BoxM(response, peanuts$variety))
```

Box's M Test

Chi-Squared Value = 10.56304 , df = 12 and p-value: 0.567

- Neither of these P-values is low enough to worry about. (Remember, the P-value here has to be *really* small to indicate a problem.)